

SLATECHECK

VERSION 0.2.0

Rights-Aware AI Governance for Creative Production

Project handbook

A compiled edition of the framework, governance model, decision and implementation specifications, data model, conformance guidance, worked scenarios, pilot materials and glossary.

Author: Ben Gunsberger

Working draft compiled 17 September 2026

Source revision 80c3537118cc

About this edition

This handbook compiles the canonical SlateCheck project documentation into a single reviewable volume. Requirement language and source ordering are preserved from the repository. Code, JSON Schemas, fixtures and machine-readable examples remain in the project repository and are referenced by the relevant chapters.

Status	Version 0.2.0 - working draft for review and controlled pilot evaluation
Pilot	The five-case usability pilot has not yet started; effectiveness and operating cost remain untested
Authority	SlateCheck supports organisational decisions and does not provide legal advice or permission by itself
Source	Repository revision 80c3537118cc, compiled 17 September 2026

Contents

Project overview	4
Framework	6
Assessment checklist specification	13
Governance model	16
Decision specification	22
Data model	30
Implementation specification	38
Conformance and validation	43
Worked scenarios	45
Interoperability and references	50
Comparative pilot protocol	51
Five-case usability smoke pilot	53
Usability smoke pilot results template	56
Migration from 0.1 to 0.2	60
Publication review	62
Glossary	64

CHAPTER 01 / README.MD

Project overview

Rights-Aware AI Governance for Creative Production

Author: Ben Gunsberger

Executive summary

SlateCheck is a governance method for deciding whether a specific use of AI is permitted within a creative business.

When someone proposes an AI activity, SlateCheck turns it into a structured case covering the exact material, the contracts and people connected to it, what the AI will do, why and where the result will be used, and the particular tool and account involved.

SlateCheck brings together the relevant facts, organisational rules and specialist decisions from Legal, Privacy, Security, Production and other responsible teams. It produces one of five clear outcomes: proceed, proceed with conditions, supply more information, obtain approval, or stop under the current rules.

Each answer applies only to the use that was assessed. It records the supporting evidence, required conditions, responsible people, validity period and changes that would require another review. Human specialists retain authority over contracts, privacy, security and approvals.

This repository currently contains a draft specification for the method: its governance model, decision rules, data structures, fictional examples, validation code and pilot plan. The planned pilot will test its usability, effectiveness and operating cost.

How SlateCheck frames a request

SlateCheck answers a practical question:

Given this material, these contributors, these agreements, this AI activity, this destination and this tool, are we permitted to proceed?

The SLATE model describes five dimensions of a proposed use:

- Source - the material, its provenance and contents
- Legal authority and contributors - the agreements, rights, people and collective terms governing it
- AI operation - what the AI system will do and retain
- Target use - the intended purpose, beneficiary, destination and audience
- Environment - the tool, account, provider, regions, access and retention settings

Slate names the complete proposed-use record; Check names the evidence-backed assessment applied to it.

SlateCheck records the evidence, authority, conditions and review triggers behind a bounded decision. It is designed for animation, VFX, games, film and other creative-production environments spanning client work, original IP, co-productions, employees, contractors and performers.

Status and intended use

Version 0.2.0 - working draft for review and controlled pilot evaluation. The [five-case usability pilot](#) has [not yet started](#). Effectiveness and operating cost remain untested in that pilot. The reference code validates examples and selected semantics; it is a specification harness.

SlateCheck supports organisational decisions. It provides no legal advice or permission by itself. Adopters configure their own policies, authorities, jurisdictions and review process with qualified legal, privacy, security, production and workforce representatives.

The canonical publication location and reuse terms remain unresolved; see the [publication review](#). This draft makes no open-licence grant. During the pilot, changes follow its evidence-led soft freeze.

Reading guide

Reader or task	Start here
Understand the framework	Framework , then worked scenarios
Ask what must be answered or approved	Assessment checklist
Set up governance and responsibilities	Governance , then conformance
Implement or exchange records	Data model , decision rules , implementation requirements and schemas
Evaluate usefulness and cost	Five-case smoke pilot , then comparative pilot

Specification set

The normative core comprises the Framework, Data Model, Decision Specification, Assessment Checklist, Governance, Implementation Specification, [Glossary](#), Conformance and the five JSON Schemas. [Framework §2](#) defines requirement language and document authority.

Supporting material is informative:

- [Scenarios](#) - fictional worked assessments and boundary cases
- [Examples](#) - five JSON examples and a preliminary checklist; the [frozen bundle](#) supports manual decision replay
- [Migration from 0.1](#) - incompatible purpose, selector and evidence changes
- [Interoperability and references](#) - conceptual relationships to ODRL 2.2, DPV 2.0, NIST AI RMF 1.0 and OAIC guidance
- [Pilot materials](#) and [comparative protocol](#) - evaluation methods and unfilled results template
- [Publication review](#) - editorial findings and remaining release decisions

Validation

With Node.js and npm, install the locked dependencies and run:

```
npm ci
npm run validate:spec
npm run check
```

`validate:spec` validates all five JSON examples and runs the test suite, also available separately as `npm test`. The bundle checks frozen evidence and verified human gate findings; replay reconstructs its decision core. Domain review, authority completeness and execution controls require separate validation. [Conformance](#) states the harness's coverage and limitations.

A distribution should retain this README, docs/, schemas/, examples/, reference/, scripts/, tests/, blank pilot templates and both package manifests so readers can follow links and repeat validation. Exclude local dependencies, credentials and live pilot or client records. Archive an identified release snapshot with its date, publisher, licence and canonical location before public distribution.

CHAPTER 02 / DOCS/FRAMEWORK.MD

Framework

Version 0.2.0 - working specification

Author: Ben Gunsberger

1. Purpose

Creative studios hold material whose permitted use depends on several overlapping sources of authority: production agreements, ownership, licences, individual engagements, consent, collective agreements, privacy obligations, workplace policy and tool terms. AI adds operations that can extract, transform, reproduce or persist characteristics of that material in new forms.

SlateCheck provides a consistent way to describe a proposed AI use, assemble the applicable evidence and return a bounded, explainable outcome. It supports human decision-making and organisational governance. It does not replace legal interpretation, consultation or approval.

The framework has four goals:

1. Give staff a practical route to an answer before material enters an AI system.
2. Let specialists maintain facts and rules in reusable modules instead of reviewing every combination from scratch.
3. Preserve the evidence, reasoning, conditions and accountability behind each answer.
4. Carry source obligations forward into transcripts, embeddings, datasets, outputs and models.

2. Normative language

The words MUST, MUST NOT, SHOULD, SHOULD NOT and MAY express requirement levels in this specification.

- MUST / MUST NOT: required for SlateCheck conformance.
- SHOULD / SHOULD NOT: expected unless a documented reason supports a different implementation.
- MAY: optional.

Core documents and schemas listed in the README are normative; worked examples, pseudocode, implementation suggestions and supporting material are informative. Uppercase requirement words define conformance obligations. The Data Model's required domain fields define record content; unknown or inapplicable values follow its completeness rules.

JSON Schemas define exchange structure, the Decision Specification defines assessment semantics, and Conformance defines level applicability. These requirements apply together. An inconsistency requires a documented resolution before relying on the affected permission; examples and reference code cannot override a requirement.

3. Unit of assessment: the proposed use

Every assessment concerns one versioned Proposed Use:

A defined set of source material, governed by identified rights and people, submitted to a stated AI operation for a stated purpose in a stated environment.

A proposed use is represented through the five-part SLATE model:

Source + Legal authority and contributors + AI operation + Target use + Environment

Changing any material fact creates a new version and triggers reassessment. Material changes include adding files or people, changing the operation, moving from one production to another, retaining a model, changing provider settings, adding an access region or distributing an output to a new audience.

4. The five SLATE dimensions

4.1 Source

Source describes the material and its provenance.

It MUST identify:

- each material item or a controlled collection with a manifest;
- category and format;
- originating production, business function or external source;
- ownership and funding context where known;
- lifecycle state, such as development, active production, delivered, archived or deleted;
- every represented contributor or a documented method for identifying them;
- personal, sensitive, confidential, culturally sensitive and security classifications;
- governing agreements and licences;
- upstream material and downstream derivatives through lineage links.

Common source categories include:

- communications: email, chat, spaces, meeting recordings and transcripts;
- production tracking: notes, client notes, versions, approvals, schedules and histories;
- supplied inputs: scripts, bibles, designs, references, audio, video, scans, motion capture and datasets;
- production outputs: artwork, animation, renders, edit media, rigs, code, tools and models;
- business and people records: contracts, HR data, performance data, finance and security logs;
- external material: licensed libraries, stock assets, public-domain material and web content;
- derived material: summaries, transcripts, labels, embeddings, indexes, synthetic outputs, adapters and model weights.

Possession, access, payment and production completion are recorded as facts. They do not independently establish permission.

4.2 Legal authority and contributors

Legal authority and contributors describes every authority and person whose terms may govern the proposed use.

It MUST support:

- production, client, co-production and distribution agreements;
- licences and assignments;
- employment and contractor agreements;
- performer, artist and other contributor agreements;
- consent records, riders and releases;
- guild, union, award and collective terms;
- privacy notices and lawful processing bases;

- organisational policies and approved exceptions;
- jurisdiction, effective date, expiry, withdrawal and supersession;
- evidence provenance and a responsible verifier.

Rights are evaluated at the most granular level required by the evidence. A dataset containing ten performers can have ten different contributor outcomes. The dataset can proceed only when the proposed use is covered for every included item and person, or when uncleared elements are removed and the manifest is updated.

An agreement summary **MUST** link to its controlled source or custodian. The summary is a queryable interpretation, not a replacement for the executed document.

4.3 AI operation

AI operation describes what the system does to the source and which durable artefacts it creates.

The minimum operation vocabulary is:

- transcribe;
- summarise;
- translate;
- classify, tag or extract;
- analyse or infer;
- search, retrieve or create embeddings;
- generate or transform content;
- evaluate or test a model;
- automate or support a decision;
- train a base model;
- fine-tune, adapt or create a LoRA;
- reinforcement learning or preference optimisation;
- distil or transfer a capability;
- create a synthetic voice, likeness or performance.

An operation **MUST** also state persistence:

- ephemeral processing;
- saved output;
- retained transcript or labels;
- embeddings or searchable index;
- curated dataset;
- model weights;
- adapter or LoRA;
- evaluation set or logs.

Technical descriptions **SHOULD** explain what the system learns or retains in production language. Labels such as “analysis” or “internal AI” are too broad for a permission decision.

4.4 Target use

Target use describes the intended outcome, beneficiary, destination and audience.

It **MUST** distinguish:

- the same task;
- the same production;
- a different production;
- a reusable studio capability;
- client delivery;
- internal operations or communications;
- research or evaluation;
- public release or external distribution.

The destination production and its ownership context are separate facts from the source production. This distinction makes cross-production reuse visible, including movement from a client-funded show to company-owned IP.

The purpose record stores production relationships, business purposes, reuse intent and distribution separately. A cross-production research experiment intended for public release records all three facts. Production relationships are derived from the source and destination IDs, including mixed same-production and different-production sets.

The purpose record **SHOULD** include the expected benefit, affected people, intended users, output audience, commercial context and impact if the system is wrong or misused.

4.5 Environment

Environment describes the complete processing boundary.

It **MUST** identify:

- product, provider, model and material version where relevant;
- deployment type: local, studio-hosted, private tenant, managed enterprise service or public service;
- account or tenant;
- provider use of inputs, outputs and telemetry, including model improvement;
- storage and processing regions;
- authorised people, teams and service operators;
- retention, deletion and backup behaviour;
- subprocessors and external integrations;
- security classification and approved data classes;
- export, monitoring and incident-response controls.

A product-level approval is insufficient when settings vary by account, region or feature. SlateCheck evaluates the actual environment used for the proposed activity.

5. Cross-cutting concepts

5.1 Positive basis

A positive basis is verified evidence sufficient for the particular gate. Rights gates require a bounded rights conclusion; inventory requires verified completeness; environment requires technical verification and approval; obligations require an executable control plan. An internal approval exercises only the authority delegated to its issuer.

SlateCheck's default is an organisational processing hold until the applicable bases are established. An authorised reviewer can record an applicable licence, ownership-based authority, legal exception or a reasoned determination

that a particular permission is unnecessary. The reviewer records the jurisdiction, exact scope, evidence and authority exercised. Consent is required where the applicable authority requires it. A non-applicability finding requires equivalent evidence and verification; silence cannot establish it.

The per-gate evidence requirements are defined in the Decision Specification. Missing facts create an information hold; known facts requiring an authorised interpretation or grant create an approval hold.

5.2 Cumulative rights

Applicable authorities accumulate. A client approval does not satisfy a performer consent requirement. An individual consent does not remove a collective minimum. Tool approval does not confer rights in the material.

5.3 Lineage and inheritance

Derived artefacts **MUST** link to their sources. Each source obligation receives a recorded applicability assessment for each relevant derivative type. The assessment identifies its authority, transformation, affected artefact, trigger and disposition: retained, released by verified interpretation, or uncertain. Uncertain applicability creates a hold on the affected downstream use until resolved. An authorised interpretation can narrow or replace an obligation only within its recorded authority.

This applies to transcripts, summaries, crops, annotations, embeddings, indexes, prompts, outputs, datasets, evaluation sets, adapters, LoRAs and model weights.

5.4 Purpose and environment binding

Permissions are bound to their recorded purpose and environment. Approval for same-show summarisation in a private tenant does not cover cross-show training or a public account.

5.5 Time and change

Evidence and decisions are time-bound. Each record **MUST** state its effective status, verification date and review or expiry trigger. Withdrawal, amendment, tool-setting changes, incidents and new destinations trigger reassessment.

5.6 Human authority

SlateCheck distinguishes an engine result from an authorised decision. Automated evaluation can confirm that a proposal matches a standing approval. It cannot create a new permission, waive a restriction or resolve an ambiguous contract.

6. Decision outcomes

SlateCheck uses five plain-language outcomes:

Outcome	Meaning	Staff action
Permitted	Current verified records positively cover the exact proposed use.	Proceed within the recorded scope.
Permitted with conditions	Positive permission exists and listed controls must be satisfied.	Verify prerequisites before processing; complete later duties at their stated time.
Hold - information required	A material fact or valid evidence record is missing, stale or inconsistent.	Supply or verify the requested information.
Hold - approval required	The facts are known and an authorised person or external party must decide.	Keep material out of the workflow until approval is recorded.
Prohibited under current rules	An applicable, verified rule expressly disallows the proposed use.	Change the proposal or obtain a valid amendment through the authorised process.

Every outcome **MUST** include:

- a plain-language answer;
- the exact assessed scope;
- blockers and conditions;
- applicable rules and evidence versions;
- decision owner and required approvers;
- validity period and reassessment triggers;
- engine or method version;
- creation time and immutable decision ID.

7. The decision record

SlateCheck can present an [assessment checklist](#) for intake and unresolved work. After formal assessment, it is a derived view; the decision and its execution prerequisites retain processing authority.

The output is a versioned evidence snapshot, not a free-floating answer. A conforming decision record **MUST** allow a reviewer to reconstruct:

1. what was proposed;
2. what facts were known;
3. which records and rule versions applied;
4. how each gate resolved;
5. who made or owned the decision;
6. which conditions remain active;
7. what change would invalidate the answer.

Previous decisions **MAY** be used as evidence of a standing approval. Precedent never broadens beyond its recorded scope.

8. Framework architecture

SlateCheck separates five services. They can begin as controlled spreadsheets and document repositories, then move to a database without changing the conceptual model.

1. Source and lineage register - productions, materials, collections, contributors and derivatives.
2. Rights and policy register - agreements, interpreted terms, consent, collective terms and internal rules.
3. Tool and environment register - approved configurations, regions, access and retention behaviour.
4. Assessment service - structured proposed uses, rule matching and unresolved questions.
5. Decision and obligation register - outcomes, approvals, conditions, expiry, deletion and audit events.

Controlled source documents remain in their systems of record. SlateCheck stores stable references, structured interpretations and verification metadata.

9. Conformance levels

All unqualified **MUST** requirements in the specification set apply at every level. A level describes how those controls are operated. Manual evidence checking, controlled registers and signed records can implement the baseline. Requirements explicitly marked Level 3 apply to automated evaluators. The normative capability matrix is in [CONFORMANCE.md](#).

The levels are cumulative: Recorded establishes controlled manual assessment and execution; Governed adds maintained workflows and standard exports; Connected adds automated evaluation and enforcement; Assured adds independent control testing and periodic measurement. The matrix specifies the required capabilities without prescribing a software stack.

10. Scope boundaries

SlateCheck is a governance and decision-record framework. It does not determine ownership, interpret law by itself, guarantee that source records are accurate, assess model quality or replace consultation with affected workers and contributors. Its value depends on verified records, clear authority and a working process for uncertainty.

An implementation **MUST** describe its jurisdictional coverage and limitations. Public examples **SHOULD** use fictional organisations, productions, people and agreements unless publication rights and privacy have been confirmed.

CHAPTER 03 / DOCS/ASSESSMENT_CHECKLIST.MD

Assessment checklist specification

Version 0.2.0 - working specification

1. Purpose

An assessment checklist answers a practical question:

For this use case, what do we need to find out, approve or complete next?

The checklist is a working view. It may be prepared manually or generated by software. The Proposed Use, decision, approval and obligation records remain authoritative.

2. Two stages

Preliminary checklist

A requester MAY start with a plain-language description before a complete Proposed Use exists. A preliminary checklist identifies known facts, assumptions and missing information across the five SLATE dimensions. Possible approvals and conditions remain labelled as potential until verified.

This stage MUST NOT issue a permission outcome or imply that every applicable authority has been found. Formal evaluation still requires a valid, confirmed Proposed Use.

Evidence-backed checklist

After assessment, the checklist is derived from the current decision snapshot. It shows the decision outcome and the unresolved information, approvals, prohibitions and conditions that govern the next action. Satisfied and non-applicable gates remain available in the linked decision rather than being repeated as checklist items.

When facts or scope change, the implementation reassesses the Proposed Use and regenerates the checklist from the successor decision. Rights conclusions remain in the decision record.

3. Required output

A saved checklist MUST identify its stage, creation time, owner, bounded use case and input references. An evidence-backed checklist MUST also identify the decision ID and version.

Each actionable item contains:

Field	Required content
Item ID and kind	A stable reference within the checklist; information, approval, condition or prohibition.
Question or action	One answerable question, bounded approval request, testable action or blocked activity.
Scope and basis	The affected gate, material/person/operation/authority as relevant, and the rule, evidence or assumption that created the item.
Owner and authority	The coordinating role and, for approvals, the authority that must decide.
Evidence and timing	What resolves the item and when it is needed.
Dependencies	Earlier items that determine applicability or resolution.

Field	Required content
Authoritative status	A link to the controlled fact, approval, obligation or successor decision that records resolution.

Implementations MAY present these fields in a document, sheet or interface. Version 0.2 defines no checklist JSON exchange schema.

4. Generation rules

1. Map the request to the five dimensions and ask scope-defining questions first.
2. Label supplied facts, assumptions and generated suggestions clearly. Only verified applicable requirements are confirmed.
3. Derive evidence-backed items from all unresolved gate findings. A higher-severity headline MUST NOT hide other blockers.
4. Keep independent authorities separate. Duplicate items may be consolidated only within one authority's recorded scope.
5. Link resolved work to controlled evidence, approval or obligation records, then reassess. A checked box or generated answer supplies no authority by itself.
6. Show prohibitions as blockers. Amendments and exceptions follow Decision Specification §6.5.

An AI-assisted implementation MUST link generated requirements to their basis, respect evidence access controls and distinguish generated wording from verified findings and exercised authority.

5. Worked example: client-show voices for an original production

Starting from a description

We want to use dialogue from an old client show to train lip sync for our own new series.

A preliminary checklist asks:

1. Which recordings, performers and source and destination productions are involved?
2. Which client, contributor and collective agreements govern them?
3. What will the model learn, generate and retain, including memorisation and deletion behaviour?
4. Who will use the model and outputs, and will either leave the destination production?
5. Which tool, tenant, regions, access groups, provider-training settings and retention rules apply?

Any approval routes remain potential until verified records establish that they apply.

Using the existing assessment

The [cross-show decision](#) produces the following compact evidence-backed checklist. The decision and its [frozen bundle](#) contain the full gate findings and evidence versions.

- Proposal: USE-HH-MM-001 version 2.0
- Decision: DECISION-HH-MM-001 version 2.0 - Hold: approval required
- Already established: the material inventory and private tenant configuration are verified within the assessed scope.

ID / kind	Question or action	Basis and scope	Owner / authority	Evidence and timing	Depends on
Q1 / information	Verify learned capability, memorisation risk, access, retention and deletion.	Operation finding; retained model	ML lead	Technical statement before final review	-
A1 / approval	Obtain authority for the recordings, training, retained model and use on <i>Moss & Moon</i> .	Source and destination findings	Business and Legal Affairs / StreamWave Kids	Scoped written authority before processing	Q1
A2 / approval	Obtain Ava's purpose-specific agreement.	Contributor finding; Ava's rider	Business and Legal Affairs / Ava Lin	Executed agreement before processing	Q1, A4
A3 / approval	Resolve Mateo's restriction and create a separate consent request if required.	Contributor finding; Mateo's rider	Business and Legal Affairs / authorised rights reviewer	Scoped finding before processing	Q1, A4
A4 / approval	Establish collective coverage separately for both performers.	Contributor finding; two collective instruments	Business and Legal Affairs / rights and labour reviewer	Individual coverage findings before contributor agreements	Q1
A5 / approval	Establish the privacy basis and contributor process for secondary use.	Privacy finding	Privacy owner / privacy authority	Verified basis and process before processing	Q1, A2-A4
A6 / approval	Obtain approval of the compensation, expiry, withdrawal and deletion plan, with assigned duties.	Obligations approval finding	Decision owner / applicable duty authorities	Approved obligation plan before processing	A1-A5

All listed items are unresolved at issue; the linked decision is their current authoritative record. Each resolution adds a link to its controlled evidence or successor decision. Q1 bounds the approval requests; agreement review can proceed in parallel. Verified responses update controlled records and produce a successor assessment.

6. Acceptance cases

A conforming checklist implementation demonstrates that:

- a description-only request produces explicit unknowns and potential approval routes without a permission outcome;
- the evidence-backed view includes every unresolved gate finding and preserves independent authorities;
- scope-defining questions precede dependent approvals;
- prohibitions remain visible with other unresolved work;
- resolution links to controlled evidence, approval or obligation records; and
- changed scope or evidence produces reassessment and a checklist derived from the successor decision.

CHAPTER 04 / DOCS/GOVERNANCE.MD

Governance model

Version 0.2.0

1. Governance objective

SlateCheck creates a shared operational system across creative, production, technology, security, privacy, People and Culture, labour relations and Legal. Its governance model assigns each fact and decision to the people best placed to own it.

The governing principle is simple: staff receive a useful answer, specialists retain authority over their domain, and every decision remains traceable to verified evidence.

2. Roles

Requester

Describes the intended use, confirms the structured proposal, follows conditions and reports changes. The requester does not interpret contracts.

Production data steward

Maintains production, material, manifest, contributor and lineage records. Confirms the exact data proposed for use.

Business and Legal Affairs

Owns contract interpretation, structured rights rules, external approvals, legal conflicts and exceptions. Legal validates summaries rather than entering every operational fact.

Casting, People and Culture, and labour relations

Maintain engagement, consent, collective-coverage and workforce-process facts. They route individual and collective changes to Legal for interpretation.

Privacy owner

Evaluates personal-data purpose, necessity, notice, consent where required, access, retention, affected-person rights and high-impact processing.

IT and security

Own the technical truth about tool environments, accounts, data regions, access, retention, logging, deletion, subprocessors and incidents.

AI governance owner

Owns the SlateCheck vocabulary, decision method, policy lifecycle, conformance, metrics and cross-functional forum. This role cannot grant rights held by another party.

Decision owner

Accepts accountability for an assessment and its next action. The owner ensures that required approvers act and conditions are tracked.

Independent assurance

Periodically tests records, decisions, controls and access. Assurance reports findings to a governance body with authority to require remediation.

3. Authority matrix

Decision area	Accountable authority	Required inputs
Contract meaning and external reuse rights	Business and Legal Affairs	Controlled agreement, proposed use, production context
Contributor consent and collective coverage	Legal with Casting / People / labour relations	Individual record, applicable collective terms, exact operation
Personal-data processing	Privacy owner or delegated qualified role	Data categories, people affected, purpose, notice and controls
Tool technical approval	IT and security	Architecture, settings, contract, regions, access, retention
Creative suitability	Creative owner	Rights-cleared inputs and outputs, intended production use
Production acceptance	Producer or production executive	Conditions, budget, schedule and client obligations
New policy or standing approval	Named governance authority	Cross-functional review and test cases
Exception	Authority named by the governing rule	Exact scope, rationale, evidence, expiry and conditions

One person may hold several roles in a small studio. The decision record still identifies which authority they exercised.

4. Operating workflows

4.1 Production onboarding

At greenlight or contract execution:

1. create the production and organisation records;
2. register controlling, funding and delivery parties;
3. link principal agreements;
4. record project IP, background IP and client-material treatment;
5. identify locations, likely contributor groups and collective instruments;
6. create initial structured rules;
7. schedule review at delivery, amendment and wrap.

4.2 Contributor onboarding

At engagement:

1. assign a controlled contributor ID;
2. link the engagement agreement and applicable collective terms;
3. record represented material categories;
4. capture purpose-specific AI consent or restriction when present;

5. record compensation, notice, consultation, expiry and withdrawal mechanisms;
6. verify the structured interpretation before use.

4.3 Material and dataset registration

Before an assessment:

1. identify the exact files or controlled collection;
2. create a versioned manifest;
3. map provenance, production and governing agreements;
4. enumerate represented people;
5. classify confidentiality and personal information;
6. verify lineage and exclusions;
7. assign a data steward.

4.4 Tool onboarding

Before an environment can satisfy a decision gate:

1. identify the actual product, model, tenant and features;
2. review vendor and data-processing terms;
3. verify provider training and human-access settings;
4. document regions, subprocessors and integrations;
5. test access, retention, export and deletion controls;
6. define approved data and operation classes;
7. record owner, evidence and review date.

4.5 Proposed-use assessment

The requester describes the use in plain language and confirms all five SLATE dimensions. Automated evaluation matches standing rules. The decision owner routes unresolved gates to the correct specialist or external party. Processing begins only when the outcome permits it and prerequisite conditions are complete.

4.6 Operation and closeout

The operational owner:

- records durable derivatives and lineage;
- completes and evidences conditions;
- monitors access and incidents;
- reviews the decision when scope changes;
- performs required deletion or archive actions;
- closes the use when retention ends.

5. Rule authoring and review

Every rule moves through:

1. Draft - created by a domain specialist from a controlled source.
2. Review - checked for scope, authority, vocabulary and test coverage.

3. Active - approved by the named authority and available to assessments.
4. Superseded, expired or withdrawn - unavailable as a positive basis and retained for audit.

Each active rule **MUST** include at least one positive and one boundary test. High-impact rules **SHOULD** be reviewed by a second qualified person.

Rule text uses plain language and controlled selectors. Contract quotations remain minimal and access-controlled. Interpretive uncertainty is represented as an approval requirement rather than hidden in notes.

6. Decision review and appeal

A requester or affected person can seek review when:

- facts are wrong or incomplete;
- relevant evidence was omitted;
- the proposal can be narrowed;
- an authorised exception may apply;
- the person is exercising a contractual, workplace or privacy right;
- the system explanation is unclear.

The review records the issue, reviewer, evidence considered, outcome and relationship to the original decision. Historical decisions remain unchanged. An appeal cannot be decided solely by the person whose original judgement is contested when organisational policy requires independence.

7. Change, withdrawal and incident response

Evidence owners publish changes as new versions. The system identifies linked active decisions, collections and derivatives. The decision owner then records one of:

- still valid;
- valid with revised conditions;
- paused pending review;
- revoked;
- expired and closed.

Consent withdrawal or a rights dispute immediately suspends affected active uses while the decision owner performs stop-use review. The suspension also covers affected derivatives whose obligation applicability is uncertain. A security or privacy incident follows the organisation's incident process and also triggers SlateCheck impact analysis.

8. Information security and privacy

A SlateCheck implementation **SHOULD** apply:

- least-privilege and role-based access;
- separate storage for identity, contract text and general metadata;
- pseudonymous contributor IDs in broad interfaces;
- encryption in transit and at rest;
- access and change logging;
- retention limits for requests and evidence snapshots;
- tested export and deletion controls;

- confidential explanations for restricted users;
- periodic access certification;
- secure handling of vendor credentials and integrations.

The system itself can contain sensitive contract and personal information. Its design requires a privacy and security assessment before live data is entered.

The organisation **MUST** set retention periods for decisions, evidence and audit events, including any applicable preservation hold. Immutability protects records during that period. Authorised disposal records the scope, authority and date without retaining deleted sensitive content; the archive marks any resulting limit on reconstruction. Active uses that lose a necessary evidence basis require reassessment.

9. Oversight forum

A cross-functional SlateCheck forum **SHOULD** meet on a defined cadence and after significant incidents. It reviews:

- new high-impact uses;
- recurring unresolved questions;
- tool and policy changes;
- contributor and workforce concerns;
- incidents and appeals;
- rule quality and inconsistent decisions;
- conformance and audit findings;
- publication of framework updates.

The forum includes meaningful creative, production and workforce representation alongside legal and technical functions.

10. Metrics

Metrics support improvement and must not become employee-performance scores. Useful measures include:

- assessment volume by operation and outcome;
- median time to resolve holds;
- proportion blocked by missing inventory, contributor or tool facts;
- expiring evidence and overdue conditions;
- decisions affected by record changes;
- rule test coverage and failed regression tests;
- incidents, appeals and reversals;
- number of active derivatives with complete lineage;
- user comprehension of answers and conditions.

Targets and service levels are configured locally. Speed never substitutes for the required authority.

11. Publication and research governance

A public SlateCheck specification can publish the framework, vocabulary, schemas, fictional scenarios and conformance tests. Implementations protect client names, contract terms, contributor identities, security architecture and decision records.

Research using operational data requires its own SlateCheck assessment, ethics and privacy review where applicable, a publication plan and a method for de-identification. Contributors and affected workers should have a route to challenge inaccurate characterisation of their rights or views.

12. Execution state and preflight

The current execution state is separate from the immutable decision outcome:

State	Meaning and next action
awaiting_prerequisites	The decision allows a conditional use; pre-processing duties remain open.
ready	A fresh preflight confirms scope, evidence, authority and prerequisite completion.
active	The operational owner has started the verified use.
suspended	A hold, changed material fact, expired basis, dispute, withdrawal, missed duty or relevant incident prevents continuation.
closed	Processing ended and remaining retention/deletion duties are recorded and tracked.

Immediately before each job or bounded batch, a named operator **MUST** verify the decision ID/version, actual manifest and environment, current evidence, open triggers, prerequisite evidence and expiry. A ready check applies only to that job/batch and scope. Recurring workflows declare batch boundaries and a maximum recheck interval in their standing approval; any intervening trigger suspends the affected work immediately. Human preflight is supported at every level; Level 3 can enforce it through integrations.

Before client delivery, external distribution or public release, the operator **MUST** verify that the actual artefacts, recipients and channel are within the decision's distribution scope and that every `before_release` duty is satisfied or validly waived. Record this release check as an execution event. A changed audience or artefact requires reassessment; a processing preflight alone cannot clear release duties.

Condition evidence or a waiver must be verified by the duty's authority. An operational owner cannot waive an external party's right. Failed or expired conditions suspend processing; future deletion duties can remain open until due. Resume requires a new documented preflight and a successor decision whenever assessed facts or scope changed. The owner records suspension, restart and closeout events without altering the historical decision.

13. Standing approvals, cohorts and operating cost

Standing approvals define source classes, verified cohorts, operations, environment configurations, purposes, distribution boundaries, exclusions, expiry and recheck cadence. New items or people enter only after the steward verifies membership and individual exceptions. Uncertain membership routes to a hold. Cohort membership and exception checks are versioned evidence.

The governance owner assigns response service levels and escalation owners for each hold type. The pilot measures specialist minutes per decision, initial record preparation, ongoing maintenance, repeated-use savings, unnecessary holds, bypass attempts and user abandonment alongside missed restrictions. Record exclusions and declined uses so the evaluation includes work that never reaches a permitted outcome. The comparative pilot protocol in [PILOT_EVALUATION.md](#) defines the evidence needed before broader rollout.

CHAPTER 05 / DOCS/DECISION_SPECIFICATION.MD

Decision specification

Version 0.2.0

1. Objective

The decision process converts a confirmed Proposed Use and current controlled records into an explainable outcome. It is deterministic for identical inputs, evidence versions and rules. Human authority remains responsible for new interpretations, approvals, exceptions and conflicts.

The process answers two separate questions:

1. Does a standing, verified basis cover this exact proposal?
2. If it does not, what information or authority is required next?

2. Inputs

An assessment requires:

- one valid Proposed Use document;
- current source, production, contributor, agreement and environment records;
- active policy rules for the relevant organisation and jurisdictions;
- the identity and authority of the requester;
- an evaluation time;
- an evaluator version.

The evaluator **MUST** preserve the input and evidence versions used. A live link to a changeable record is insufficient for audit.

3. Assessment gates

Each proposal is evaluated through eight gates. Every gate produces `satisfied`, `conditional`, `missing_information`, `approval_required`, `prohibited` or `not_applicable`.

Gate	Question
Inventory	Do we know the exact material and every represented contributor?
Source authority	Is there a positive basis to use these sources in this way?
Contributor rights	Are individual, employment and collective requirements satisfied for every person?
Privacy and people impact	Is the personal-data processing justified, proportionate, transparent and appropriately controlled?
Operation	Does authority cover what the system learns, produces and retains?
Purpose and destination	Does authority cover the beneficiary, production, audience and reuse boundary?
Environment	Is the exact processing environment approved for this material and activity?

Gate	Question
Obligations	Can all required notices, payments, consultations, reviews, deletion and access controls be completed and evidenced?

Organisations MAY add gates for export control, Indigenous data governance, child safety, regulated content, national security or other domain needs. Added gates use the same status vocabulary.

4. Resolution procedure

Step 1 - Validate and normalise

Validate the request against the Proposed Use schema. Resolve controlled IDs, normalise vocabulary values and preserve explicit unknowns. Reject malformed records as a system error rather than a rights outcome.

Step 2 - Expand the source set

Expand every material collection to its manifest. Traverse upstream lineage required to identify governing sources, agreements and represented people. Record the version of every traversed edge.

Step 3 - Build the authority set

Collect all potentially applicable:

- production and client agreements;
- ownership, licence and supplier records;
- individual engagements, consents and riders;
- guild, union, award and collective instruments;
- privacy notices and processing bases;
- internal policies and exceptions;
- environment approvals and vendor terms;
- prior standing decisions.

Expired, superseded or withdrawn records remain in the audit trail and do not provide a current positive basis.

Step 4 - Match rules

Rules use the set predicates and relationship binding in §12. All populated predicates and retention constraints are combined with AND. Omitted selectors are unrestricted within the rule's verified authority; an empty `values` list is invalid. Free text explains a rule and cannot satisfy a selector. A matched rule is a candidate for the gates named in its `gates` field. Its effect applies only within its authority and matched scope.

Rule matching MUST consider:

- the whole material set;
- each represented person individually;
- every relevant jurisdiction;
- both source and destination production;
- durable artefacts created by the operation;
- actual environment settings.

Step 5 - Evaluate each gate

For each gate:

1. establish the gate-specific basis and the complete set of applicable authorities described in §13;
2. apply prohibitions and mandatory requirements from all applicable authorities;
3. determine whether evidence is current and internally consistent;
4. identify conditions that can be completed operationally;
5. identify decisions reserved for a named authority.

A positive basis must cover the full scope. Partial coverage creates a hold unless the proposal can be partitioned into separately assessed collections.

Step 6 - Aggregate the outcome

The primary outcome follows the highest-severity gate result:

1. `prohibited_under_current_rules`
2. `hold_approval_required`
3. `hold_information_required`
4. `permitted_with_conditions`
5. `permitted`

All discovered blockers and missing facts remain visible even when a higher-severity result determines the headline.

Step 7 - Create the decision snapshot

Write an immutable decision containing the confirmed proposal, gate results, explanation, applicable rule IDs, evidence versions, conditions, owner, validity and reevaluation triggers.

When presenting an [assessment checklist](#), derive its evidence-backed items from this snapshot and include every unresolved gate finding. Satisfied and non-applicable findings remain in the decision. Preliminary questions may precede Step 1; verified answers feed a new assessment and successor decision.

Step 8 - Register downstream artefacts

When processing proceeds, create material and lineage records for durable outputs. Attach the decision ID and inherited controls. Conditions remain open until completion evidence is recorded.

5. Evaluation contract

Implementations follow §4 in order and preserve the evaluation time, expanded scope, evidence versions, matched rules and gate findings needed to reproduce the result.

Each gate evaluation **MUST** require the gate-specific basis when the gate is applicable and an evidenced non-applicability finding otherwise. Silence resolves to `approval_required` when the relevant facts are known and an authority must interpret or grant rights. Missing, stale or contradictory facts resolve to `missing_information`.

6. Rules and conflict handling

6.1 Rule effects

SlateCheck supports five effects:

- `permit` - supplies a positive basis for a bounded use;
- `prohibit` - expressly disallows a bounded use;
- `require_condition` - attaches an operational duty;
- `require_information` - identifies evidence required before resolution;

- `require_approval` - reserves the decision for an identified authority.

6.2 Cumulative application

Rules from independent authorities apply cumulatively. For example, a client permission, performer consent, collective requirement, privacy basis and environment approval may all be required.

6.3 Specificity

Specific rules can refine general organisational rules when the specific rule has authority to do so. A valid, signed project exception can narrow a general internal policy for one named dataset and environment. It cannot override an independent client's rights or a contributor's agreement unless its authority expressly includes them.

6.4 Conflicts

The evaluator **MUST NOT** invent legal precedence. It may apply precedence explicitly recorded in an authorised policy, supersession record or agreement interpretation. An unresolved conflict between applicable records produces `hold_approval_required` and identifies the decision owner.

An express prohibition determines `prohibited_under_current_rules` when the evidence is current and no authorised exception or superseding rule applies.

6.5 Overrides and exceptions

An override is a new controlled record, not an edit to historical evidence. It **MUST** identify:

- the rule or decision being overridden;
- the authority and evidence supporting the override;
- exact scope;
- conditions;
- effective and expiry dates;
- approver and rationale.

7. Contributor-set evaluation

Each represented person is evaluated separately against individual and collective records. The collection gate aggregates those results:

```
if any person is prohibited: prohibited
else if any person requires approval: hold for approval
else if any person's evidence is missing: hold for information
else if any person has conditions: permitted with combined conditions
else: satisfied
```

Group assumptions such as “all employees signed the same contract” require a verified cohort record and an exception check. Named individual riders always remain discoverable.

8. Derived artefacts

The requested operation must declare each expected durable artefact. The evaluator treats those artefacts as prospective derivatives and tests whether the positive basis covers:

- creation;
- retention duration;
- access;
- use on the destination production;

- further training or adaptation;
- distribution;
- deletion or withdrawal response.

“Training” is incomplete without naming whether the result is a base model, fine-tune, LoRA, embedding index, classifier, evaluation set or another retained capability.

9. Decision validity

A decision remains valid only while its facts, evidence and conditions remain valid. Reassessment triggers include:

- source collection or contributor-set changes;
- agreement amendment, expiry, withdrawal or dispute;
- destination or audience change;
- operation, output or retention change;
- tool, model, account, provider term, region or access change;
- new legal or collective requirements;
- missed condition;
- security or privacy incident;
- scheduled review date.

Affected decisions SHOULD be discovered automatically through reverse links. Reassessment creates a new decision and preserves the old record.

10. Explanation requirements

The answer shown to a requester MUST:

- lead with the outcome in plain language;
- describe the exact bounded use;
- distinguish known restrictions from missing information;
- name the next responsible role;
- state each condition as a testable action;
- show validity and change triggers;
- provide traceable evidence references appropriate to the requester's access.

It MUST NOT expose confidential contract text or personal information to unauthorised users. A redacted explanation can refer to a verified rule ID and owner.

11. Required evaluator tests

A conforming implementation MUST document positive and boundary cases. A Level 3 evaluator MUST automate tests for:

1. same-show summarisation with a standing approval;
2. cross-show reuse where client permission is absent;
3. a mixed contributor set with one restrictive rider;
4. an expired consent;

5. provider training enabled in an otherwise approved tool;
6. a changed processing region;
7. a missing material manifest;
8. conflicting contract interpretations;
9. a valid specific exception;
10. a derivative whose source permission is withdrawn;
11. automated employment ranking prohibited by policy;
12. a rule update that invalidates a previously permitted decision.

Tests **MUST** assert the outcome, gate results, evidence set, conditions and responsible owner.

12. Selector and normalisation semantics

A selector is { "operator": "all_in", "values": ["AU"] }, with a non-empty, unique controlled value set. Actual facts are sets. The operators mean:

Operator	True when
any_of	At least one known actual value is in the selector set.
all_in	Every actual value is known and belongs to the selector set.
contains_all	Every selector value occurs among the known actual values.
equals_set	Both known sets have exactly the same members.

A confirmed empty actual set matches none of these operators. Missing facts, null and unknown are unknown. With partial knowledge, a known intersection establishes `any_of`; all required known members establish `contains_all`; a known disallowed member disproves `all_in` and `equals_set`. Other unresolved comparisons remain unknown. In an AND expression, false dominates unknown, and unknown dominates true. Unknown never supplies a positive basis. An unknown potentially applicable requirement is retained as an information blocker.

Whole-proposal selectors on a permit rule **MUST** use `all_in` or `equals_set`: operations, persistence, destination productions, production relationships, business purposes, reuse intent, distribution, audiences, environment IDs, deployments, provider training, processing regions, access groups and jurisdictions. This prevents an Australia-only permission from covering an Australia-and-US environment. `any_of` can detect a prohibited member or trigger an approval requirement. An omitted whole-proposal selector is unrestricted only within the verified source authority; rule authors must substantiate that breadth.

Source selectors bind to a single material/person relationship. Each expanded row contains one material, its production, one relevant person (or an explicitly verified absence), and agreements linked to that material/person. A rule requiring Show A and Person B matches only where both belong to that row. Values from unrelated assets **MUST NOT** be joined to manufacture a match. Contributor-specific agreement, role and collective selectors use that person's linked records. Jurisdictions are the full applicable set for the proposal and are not inferred from a person's location alone.

A collection-wide prohibition is established by any verified matching row. Positive coverage is evaluated for every required material/person/authority combination; one permitted row cannot clear the collection. Verified cohort records can supply individual coverage only with a current membership manifest and exception check.

`constraints` contains typed `lte_days` predicates for source, output, log, backup and derivative retention. Numeric durations are non-negative days; `session_only` is zero for comparison, `indefinite` exceeds every finite limit, and unknown or unevidenced `not_applicable` is unresolved. `derivatives.retention` evaluates all

retained derivatives. A verified absence of retained derivatives satisfies that constraint. Retention requirements also need an operational duty and evidence mechanism.

Production relationships are derived from all source/destination pairs. Equal IDs yield `same_production`; different IDs yield `different_production`; a mixed set retains both. A verified absence of a source or destination production yields `no_production`; unknown IDs yield `unknown`. Business purpose, reuse intent and distribution never replace this derivation. Unknown vocabulary values and contradictory declared/derived facts are validation errors or explicit information holds; implementations **MUST NOT** silently coerce them into permission.

Effective intervals are start-inclusive and end-exclusive in UTC. At a review deadline, a record needs reverification before supplying a positive basis. Every evaluator records its evaluation time and policy-profile version. Inactive rules are excluded from active matching; uncertain applicability of potentially governing evidence remains a blocker and must not disappear during filtering.

13. Gate bases and authority coverage

Gate	Required basis	Accountable verifier
Inventory	Verified manifest, contributor mapping, exclusions and relevant lineage	Data steward
Source authority	Bounded ownership/licence/contract/legal conclusion for each relevant source authority	Authorised rights reviewer
Contributor rights	Individual and collective coverage, including an evidenced determination where a permission is unnecessary	Authorised rights and labour reviewer
Privacy and people impact	Purpose-specific processing basis and required impact/people controls	Privacy owner
Operation	Rights coverage plus verified technical account of processing and persistence	Rights reviewer and technical owner
Purpose and destination	Rights coverage for beneficiaries, production boundaries, reuse and distribution	Authorised rights reviewer
Environment	Current verified configuration and approval for the data/operation classes	IT/security and relevant privacy authority
Obligations	Assigned, feasible, testable duties, prerequisites and consequences	Operational owner and duty authority

Each satisfied or conditional gate records `basis` entries with kind, evidence, verifier and scope. A `not_applicable` gate records a `non_applicability` basis with equivalent evidence. Examples include a verified non-personal source or a qualified conclusion that a specific rights permission is unnecessary. A model's quality or an internal risk acceptance cannot supply a missing external right.

Before claiming complete coverage, the reviewer **MUST** attest that the applicable authority set is complete for the declared jurisdictional and source scope. Missing authority records create an information hold. For each rights question, preserve the material, person where relevant, operation, authority, supporting evidence and finding. Independent authority requirements accumulate; permission from one authority can cover multiple items only within its scope. In automated evaluation, these tuples form an explicit coverage proof, and unresolved tuples prevent permission.

An unresolved conflict in meaning is routed to the appropriate authority. A verified independent prohibition still determines the headline outcome; a conflict about whether that prohibition governs must first be resolved as an approval hold. Rule specificity alone supplies no override authority.

14. Validation and replay boundaries

Structural validation checks the JSON shapes, core gate cardinality and effect-specific fields. Semantic validation checks outcome aggregation, evidence closure, versions, time, basis kinds, approvals, conditions and purpose consistency. Domain validation establishes whether facts, rights interpretations, authority coverage and operational controls are substantively adequate. All three are required before operational reliance.

The repository's reference helpers implement set/retention matching, bound source rows, aggregation, preflight readiness and selected semantic invariants. They are a specification test harness. They do not implement a complete Level 3 evaluator.

`cross-show-lipsync.bundle.json` freezes the manual/hybrid assessment. Replay reconstructs gate results, the primary outcome and the applicable rule set from independently stored, verified gate findings. Those findings are human interpretation inputs. The replay validates their consistency with the exported decision and frozen evidence. It does not independently derive legal interpretations from contract text or certify authority completeness.

CHAPTER 06 / DOCS/DATA_MODEL.MD

Data model

Version 0.2.0

1. Model shape

SlateCheck is a linked record system. It avoids one giant permission matrix and avoids precomputing every possible combination. Specialist teams maintain authoritative facts and reusable rules; an assessment joins the relevant records for one proposed use.

The model has four layers:

Evidence layer	productions, materials, people, agreements, tools
Policy layer	structured permissions, prohibitions, duties and review triggers
Transaction layer	proposed uses and requested approvals
Audit layer	decisions, evidence snapshots, conditions and lifecycle events

Every record has a stable ID, version, responsible owner and timestamps; registry records also carry lifecycle status. Relationships use IDs so the storage system can change without breaking the model.

2. Common record envelope

Registry records, including policy rules, MUST contain the following envelope. `effectiveUntil` and `reviewAt` are optional when no such date applies. Proposed uses and immutable decisions are transaction records with their own explicit metadata in the published schemas; the registry envelope does not apply to them. `recordType` is `policy_rule` for a rule.

Every registry record MUST contain:

Field	Meaning
<code>id</code>	Permanent, unique identifier
<code>recordType</code>	Controlled record category
<code>version</code>	Immutable version identifier
<code>status</code>	Draft, active, superseded, expired, withdrawn or deleted
<code>name</code>	Human-readable label
<code>ownerId</code>	Accountable record owner
<code>sourceSystem</code>	Authoritative system or repository
<code>sourceReference</code>	Controlled link or document reference
<code>effectiveFrom</code>	Start of applicability
<code>effectiveUntil</code>	End of applicability, if any
<code>verifiedAt</code>	Last substantive verification
<code>verifiedBy</code>	Qualified verifier
<code>reviewAt</code>	Scheduled review date, if any
<code>changeTriggers</code>	Events that require reverification
<code>createdAt, updatedAt</code>	Audit timestamps

Field	Meaning
classification	Access and sensitivity classification

Records containing contract interpretations or sensitive personal information **MUST** have access controls independent from the general staff query interface.

Required domain fields in §3 **MUST** be accounted for, with a verified value, explicit unknown or evidenced non-applicability. Dates for future lifecycle events may remain unknown. Incomplete intake stays provisional; a record used as a positive basis **MUST** have verified facts for its assessed scope. Where the exchange schema has no suitable unknown representation, retain the request as preliminary intake rather than inventing an ID or a completed fact.

3. Core records

3.1 Organisation

Represents the studio, client, production partner, vendor, licensor, union, guild or other legal entity.

Required domain fields:

- legal and trading names;
- organisation role or roles;
- relevant jurisdictions;
- parent or affiliate relationships;
- contact or responsible role reference.

3.2 Production

Represents a show, film, game, development project or internal production.

Required domain fields:

- production type: client service, company-owned IP, co-production, licensed IP, pitch/development or internal media;
- owning and controlling parties;
- funding parties;
- operating locations;
- lifecycle state;
- principal agreement IDs;
- delivery, wrap and archive dates;
- background-IP and project-IP treatment where relevant.

The production type is descriptive. Permission comes from linked rules and evidence rather than the label alone.

3.3 Person or contributor

Represents a performer, artist, employee, contractor, participant or other identifiable person.

Required domain fields:

- internal pseudonymous or controlled person ID;
- relationship and role categories;
- applicable jurisdictions;
- engagement agreement IDs;

- collective-instrument IDs;
- consent-record IDs;
- represented-material links;
- withdrawal or dispute status.

General interfaces SHOULD display the minimum personal information needed for the decision. Contract details and identity data remain in restricted source systems.

3.4 Agreement

Represents an executed contract, licence, rider, release, consent, collective instrument, privacy notice, policy or formal exception.

Required domain fields:

- agreement type;
- parties and their roles;
- covered productions, materials and people;
- jurisdictions;
- execution and effective dates;
- source document reference;
- supersession relationships;
- interpretation owner;
- structured term IDs.

Collective terms and individual terms are separate agreement records even when they apply to the same person.

3.5 Policy rule

Represents one queryable interpretation of an authoritative source. A rule contains:

- effect: permit, prohibit, require condition, require information or require approval;
- authority and evidence references;
- selectors that define when it applies;
- conditions, duties or approval authorities;
- effective dates and jurisdictions;
- interpretation confidence and verifier;
- supersession and exception relationships.

One agreement usually creates several rules. Rules SHOULD remain atomic: one effect, one intelligible proposition and a bounded scope.

3.6 Material item

Represents a file, message, recording, shot version, design, script, rig, code repository, note, dataset, model or other source or derivative.

Required domain fields:

- material category, format and storage reference;
- originating production or business function;
- creator, supplier and represented-contributor IDs;

- ownership and licence evidence IDs;
- personal and sensitive-data classifications;
- confidentiality and security classification;
- creation, delivery and archive dates;
- lineage links;
- integrity identifier such as a manifest entry, content hash or immutable version.

3.7 Material collection

Represents a controlled set of material items used as one dataset, corpus, pack or batch.

It **MUST** contain or reference a versioned manifest. Its rights position is the intersection of the requirements applying to all included items and represented people. Removing an item produces a new collection version.

3.8 Tool environment

Represents an approved configuration, not a product brand in the abstract.

Required domain fields:

- product, provider, model and version;
- deployment and tenant type;
- contract and data-processing agreement IDs;
- approved material and operation classes;
- provider-training and human-review settings;
- processing, storage and support-access regions;
- authorised groups;
- source, output, log and backup retention;
- deletion mechanism and evidence capability;
- subprocessors, plugins and connected services;
- security approval and review dates.

3.9 Proposed use

Represents the request being assessed. It contains the five SLATE dimensions and references controlled records wherever possible. The normative machine shape is defined in `../schemas/slatecheck-proposed-use.schema.json`.

The request **MUST** preserve unknown values explicitly. A user-friendly front end can collect plain language first, then require confirmation of the structured proposal.

Use the schema's unknown or nullable values where available. Omitted optional facts remain unknown unless verified evidence establishes non-applicability. Empty contributor lists establish absence only with a verified inventory; empty production-ID lists mean verified absence of a production, while omission means unknown. Normalisation **MUST** preserve this distinction before rule matching.

3.10 Decision

Represents the immutable result for one proposed-use version. It includes:

- outcome and plain-language answer;
- evaluated gates;

- blockers and active conditions;
- approvals;
- applicable rule IDs;
- evidence versions or hashes;
- validity and reassessment triggers;
- responsible owner;
- evaluation method and version.

The normative machine shape is defined in `../schemas/slatecheck-decision.schema.json`.

3.11 Obligation and event

Represents a required action and evidence of its lifecycle. Examples include consultation, payment, source deletion, model deletion, notice, attribution, human review and access removal.

An obligation has an owner, due date, status, evidence reference and consequence of non-completion. Completion does not alter the original decision; it adds a linked event.

The [assessment checklist](#) is a working view over these records. Core record and authority status remain with the applicable fact, approval, obligation and decision records. A saved evidence-backed checklist references the decision version from which it was generated.

4. Required relationships

At minimum, a SlateCheck implementation **MUST** support these relationships:

```

Production --governed by--> Agreement
Person     --party to-----> Agreement
Person     --covered by----> Collective instrument
Material   --originated in-> Production
Material   --represents----> Person
Material   --governed by---> Agreement
Collection --contains-----> Material
Derivative --derived from--> Material or Collection
Rule       --supported by--> Agreement or Policy
Rule       --applies to----> SlateCheck selectors
Proposed use--uses-----> Material or Collection
Proposed use--runs in-----> Tool environment
Decision   --assesses-----> Proposed use version
Decision   --relies on----> Record versions
Decision   --creates-----> Obligation
Event      --satisfies----> Obligation

```

Relationships **MAY** be stored in relational join tables, graph edges, document references or controlled spreadsheet tabs. Their semantics remain the same.

5. Source-to-derivative lineage

Lineage answers three questions:

1. Which exact sources contributed to this artefact?
2. Which people and agreements are represented through those sources?
3. Which active restrictions, conditions and expiry events must be reconsidered?

Each lineage edge records:

- parent and child IDs and versions;
- transformation or operation;
- date and responsible process;

- whether content, features, labels or capability were transferred;
- contribution granularity when known;
- governing decision ID.

A model record **SHOULD** describe whether source influence can be isolated or removed. When removal cannot be reliably demonstrated, the model inherits the strictest relevant deletion or stop-use trigger pending human review.

6. Rule selectors

Policy rules are modular because they match facts across standard selector groups:

Selector group	Examples
Source	production, material category, supplier, classification, lifecycle state
Legal authority/contributor	agreement, party, contributor role, collective instrument, consent state, jurisdiction
AI operation	summarisation, embeddings, generation, training, synthetic performance, automated decision
Target use	same production, cross-production, reusable capability, public release, audience
Environment	tool, tenant, deployment, region, access group, provider training, retention
Time	effective period, expiry, production stage

Selectors may use controlled categories and explicit IDs. Free-text selectors **MUST NOT** drive an automated permission outcome.

7. Data stewardship

Record family	Supplies facts	Verifies interpretation	Typical change event
Productions and ownership	Production operations	Business and legal affairs	Greenlight, amendment, delivery, wrap
Materials and lineage	Department or data steward	Production technology	Ingest, edit, collection change, archive
Contributor identity and engagement	Casting / People and Culture	Business and legal affairs	Engagement, rider, withdrawal, role change
Collective terms	People and Culture / labour relations	Qualified legal adviser	New agreement, jurisdiction or term
Privacy basis and notices	Privacy owner	Privacy or legal adviser	New data class, purpose or notice
Tool environment	IT / security / procurement	Security, privacy and AI governance	Vendor, feature, region, setting or contract change
Policy rules	Domain owner	Named authority	Interpretation, policy or source change
Proposed use	Requester and project owner	Relevant data steward	Any scope change
Decision and conditions	Decision owner	Required approvers	Approval, expiry, appeal or evidence change

The person entering a fact and the person authorising its interpretation **SHOULD** be separately identified.

8. Storage and access pattern

A first implementation can use a controlled Google Sheet with one tab per record family, stable IDs and protected ranges. Contract documents, identity data and sensitive evidence remain in access-controlled repositories; the sheet stores references and approved structured summaries.

A production implementation SHOULD provide:

- field-level access controls for personal and contractual data;
- append-only decision and event history;
- immutable version snapshots;
- validation against controlled vocabularies;
- automated reminders for expiry and reverification;
- impact queries from changed evidence to affected decisions and derivatives;
- export in a documented machine-readable form;
- tested backup, retention and deletion procedures.

9. Minimum viable dataset

A useful pilot does not require every historic asset. It needs enough connected records to exercise the selected real cases: representative productions and materials, contributors with differing terms, verified agreement summaries, materially different tool environments, the rules needed for those cases, expected outcome states and at least one derivative lineage chain.

The pilot SHOULD measure how often answers are blocked by missing facts. That reveals which records need operational ownership before further automation.

10. Obligation applicability and execution records

An obligation applicability record MUST identify the source obligation and its version, source authority, parent and child versions, transformation, affected artefact category, triggering event, disposition (`retained`, `released_by_interpretation`, `uncertain`), verifier and evidence. It MUST state the downstream consequence: restrict access, stop use, delete, quarantine, retire or retrain. A release cites a qualified interpretation covering the exact transformation and scope. Uncertainty remains actionable and visible.

Source-file deletion, output deletion, index removal and model retirement are distinct obligations. Removing a file from a training directory does not evidence removal of its influence from model weights. Where isolation is unverified, record quarantine, retirement or retraining as the pending response, with an owner and deadline.

An execution event records the decision ID/version, operator, time, actual manifest and environment versions, validity check, trigger check, prerequisite evidence and resulting execution state. Events append to history; they do not change issued decisions. Obligation completion and waivers similarly append evidence-backed events.

11. Frozen evidence bundles

The assessment-bundle schema packages a proposed use, decision, rules and controlled evidence records. Every referenced evidence and rule version MUST resolve to frozen content, including the evaluation profile, authority dependencies, collective terms and technical approvals. Records may point to restricted documents, but the archived snapshot must retain the verified facts and interpretations actually used. A hash alone cannot reconstruct unavailable content.

The evidence-record schema validates registry metadata and dependency references. Its `facts` object carries the record-family content described in this document. Each implementation **MUST** validate that domain content against its declared profile; passing the envelope schema alone establishes no rights.

The example bundle uses fictional controlled facts and verified gate findings. Its evidence records are projections of the facts used in the assessment, rather than complete operational registry exports; omitted domain facts remain outside the example's claims. It supports replay of a manual decision's gate aggregation. Automated implementations additionally preserve expanded source/person bindings, rule-match results and authority coverage proofs.

Bundle snapshots include a `sha256`: digest of each archived record using the `recordDigest` serialization: recursively sort object keys in UTF-16 lexicographic order, rebuild objects, then apply JavaScript `JSON.stringify` and hash the UTF-8 bytes. Integer-index keys follow JavaScript's numeric enumeration order; arrays retain their order, numbers use JavaScript serialization, and strings retain their Unicode form. Other runtimes **MUST** reproduce these bytes. Readers verify versions and digests; archive controls and verifier authority establish provenance.

12. Exchange versions and extensions

Version 0.2 uses JSON Schema Draft 2020-12 and versioned schema URNs. Load the five schemas together to resolve their references. A bundle declares `specVersion`; standalone records **MUST** travel with their schema URN in exchange metadata. A record's `version` identifies its content revision, independently of the specification version. Unsupported schema or profile versions require review before operational reliance.

Most exchange objects reject undeclared fields. Extensions **MUST** use declared schema extension points. Profiles **MAY** define namespaced domain fields inside evidence `facts` and extension gates through `organisation_defined` and `organisationGateName`. A versioned policy profile **MUST** document its vocabularies, domain validation, added gates and basis requirements. Other extensions require separately identified schemas or linked records. Extensions **MUST** preserve core meanings; unknown values or unsupported extensions affecting permission remain unresolved.

Implementation specification

Version 0.2.0

1. Purpose and system boundary

This document defines what a SlateCheck implementation must do without prescribing an interface or technology stack. A conforming implementation can be a controlled spreadsheet workflow, database-backed service, production-management integration or another system that satisfies the framework and audit requirements.

The implementation receives a proposed use, retrieves controlled facts and rules, produces an explainable result, routes unresolved work and preserves an immutable record. Source contracts, media and identity records can remain in their existing repositories.

2. Actors

The implementation supports the requester, data steward, rights editor, verifier, environment owner, decision owner, approver, auditor and administrator roles. [Governance §2-3](#) defines responsibilities and authority. One person may hold several roles; permissions attach to roles and records.

3. Functional requirements

FR-01 - Capture a proposed use

The system **MUST** capture and version all five SLATE dimensions. It **MUST** preserve unknown values and **MUST** require confirmation before evaluation.

FR-02 - Identify exact material

The system **MUST** accept individual material records or a versioned collection manifest. It **MUST** prevent a permitted outcome when the required inventory or represented-person set is unknown.

FR-03 - Link controlled evidence

The system **MUST** link productions, materials, people, agreements, policy rules and tool environments through stable IDs. It **MUST** record evidence version and status in every decision snapshot.

FR-04 - Evaluate contributors individually

The system **MUST** evaluate each represented person against their individual and collective records. It **MUST** support narrowing a collection and reassessing the new version.

FR-05 - Resolve modular rules

The system **MUST** use the structured selector semantics and preserve rule effect, authority, explanation and requirements. Levels 1 and 2 may record human-verified matching and findings. A Level 3 evaluator **MUST** automate deterministic matching and authority coverage proofs.

FR-06 - Require a positive basis

The system **MUST** require the gate-specific basis in the Decision Specification, including evidenced non-applicability where appropriate. Silence and ambiguity **MUST** produce a hold.

FR-07 - Return a standard outcome

The system **MUST** use the five outcomes defined in the framework. It **MUST** show all blockers and unknowns even when one determines the headline result.

FR-08 - Route unresolved work

Every missing item, approval and condition **MUST** have a responsible role, timing and evidence requirement. The system **SHOULD** support assignment and reminders.

FR-09 - Record human authority

The system **MUST** distinguish a rules-engine result from an authorised human decision. Approvals and exceptions **MUST** record the person, authority exercised, date, scope and evidence.

FR-10 - Preserve immutable decisions

Once issued, a decision **MUST** be immutable. Corrections or reassessments create linked successor decisions.

FR-11 - Track conditions

The system **MUST** track conditions as open, satisfied, waived by an authorised person, failed or expired. Completion **MUST** reference evidence.

FR-12 - Maintain lineage

The system **MUST** link derivatives to relevant sources and their creation decision, record obligation applicability and identify affected uses when a source changes. Levels 1 and 2 may operate these controls manually; Level 3 automates them.

FR-13 - Manage validity

The system **MUST** represent effective dates, review dates, expiry, withdrawal and supersession. It **SHOULD** automatically flag affected active decisions.

FR-14 - Explain safely

The system **MUST** provide a plain-language answer while respecting access controls on contract text and personal information. Redacted users receive the verified rule ID, effect, owner and required action.

FR-15 - Export and test

Level 2 and above **MUST** support the published machine-readable exports. Every level **MUST** preserve reconstructable records and pass documented rule regression cases before activation. Level 3 **MUST** automate those cases.

FR-16 - Generate an assessment checklist

The system **MUST** support a use-case-specific [assessment checklist](#), manually or through software. Preliminary items distinguish facts, assumptions and potential requirements. Evidence-backed items are derived from the current decision, include every unresolved gate finding and link resolution to controlled records. The checklist **MUST NOT** become an independent permission or execution-readiness record.

4. Quality requirements

Security

- Role-based access applies to record type, field and source reference where required.
- Authentication uses the organisation's managed identity service.
- Administrative and approval actions require strong authentication.

- Changes, reads of restricted records and exports are logged.
- Secrets and vendor credentials remain outside SlateCheck records.

Privacy

- Interfaces expose the minimum identity and contract data required for each role.
- Broad interfaces use contributor IDs or approved display names.
- Search and analytics respect record-level permissions.
- Retention and deletion rules cover SlateCheck's own logs, requests and snapshots.
- Production use receives a privacy impact assessment appropriate to its jurisdictions and data classes.

Reliability and integrity

- Stable IDs never point to a different logical record.
- Versions are immutable after verification.
- Evaluation is repeatable from the stored snapshot.
- Backups, restore and disaster recovery are tested.
- Clock, timezone and date semantics are explicit.

Usability

- Requesters can describe intent in production language.
- Structured interpretation remains visible and editable before confirmation.
- Outcomes lead with the action the user can take.
- Conditions are concrete, assigned and time-bound.
- The system distinguishes missing information, required approval and prohibition.

Interoperability

- Records use documented controlled vocabularies and stable IDs.
- Level 2 and above import and export JSON matching the published schemas and semantic invariants.
- Source references can point to contract, production, identity and asset systems without copying protected contents.
- Extensions follow the version and extension rules in [Data Model §12](#).

5. Reference service boundaries

A mature implementation can expose these logical services:

Service	Responsibility
Registry	Controlled productions, people, materials, agreements and environments
Policy	Versioned rules, tests, authority and vocabulary
Assessment	Validation, expansion, rule matching and gate evaluation
Workflow	Assignments, approvals, conditions, reminders and appeals
Lineage	Source, derivative and impact relationships
Audit	Immutable decisions, evidence snapshots, events and exports

Services may share one datastore in a pilot. Their responsibilities remain separate.

6. Controlled-sheet pilot

A controlled workbook can test the model before software development. One possible structure uses these tabs:

Tab	Key columns
Productions	ID, type, owners, funders, locations, state, agreement IDs
Materials	ID, version, category, production, represented people, classifications, source reference
Collections	ID, version, manifest reference, material IDs, steward, status
People	Controlled ID, role, jurisdiction, agreement IDs, collective IDs, consent status
Agreements	ID, type, parties, scope, effective dates, controlled document reference, verifier
Rules	ID, version, effect, selectors, requirements, authority, dates, verifier
Environments	ID, product, tenant, regions, access, provider training, retention, approval
Requests	Proposed Use fields and confirmation state
Decisions	Outcome, scope, gates, rules, evidence snapshot, owner, validity
Obligations	Decision, condition, owner, due date, status, completion evidence
Lineage	Parent, child, operation, decision, transfer type, date
Vocabulary	Allowed values, definitions and status

Protected ranges assign editing rights by stewardship domain. The request experience can initially use a Google Form or a separate intake sheet. Decision rows become append-only after issue. Sensitive contract text and identity data stay outside the workbook.

The pilot succeeds when the team can answer representative questions consistently, discover missing source records and update a rule without rewriting every scenario.

7. Minimum pilot workflow

A requester may first receive a preliminary checklist to gather missing facts. After assessment, a decision-derived checklist routes the remaining work.

1. Requester submits and confirms a Proposed Use.
2. Data steward verifies the material collection and represented people.
3. A simple evaluator matches active rules and creates gate results.
4. The decision owner reviews the explanation and routes holds.
5. Qualified approvers record new decisions where required.
6. A decision snapshot is issued.
7. Conditions and derivatives are tracked to closure.
8. Evidence changes trigger an impact report and reassessment queue.

The pilot may calculate results manually. It still uses the same record structure and outcome rules.

8. Acceptance criteria

An initial implementation is ready for controlled internal use when it can demonstrate the following. Automated evaluation is required only at Level 3; lower levels use documented human controls. Machine-readable export criteria apply at Level 2 and above:

1. published examples validate, and every active rule has authority, verification, effective status and positive and boundary tests;
2. the same frozen snapshot produces the same gate results and resolves every cited rule and evidence version;
3. contributor and material relationships are evaluated without combining unrelated rows, and whole-scope selectors do not turn partial coverage into permission;
4. unknown, stale, expired, withdrawn or changed facts cannot provide a positive basis and trigger reassessment or suspension as applicable;
5. explicit prohibitions and every lower-severity blocker remain visible;
6. every satisfied or non-applicable gate has the required evidence, scope and verifier;
7. decisions, prerequisites, conditions and retained or released derivative obligations remain reconstructable, with owners and completion evidence;
8. reverse lineage identifies decisions and derivatives affected by a source change;
9. access, explanation, export, backup and deletion controls pass security and privacy testing;
10. preliminary and evidence-backed checklists pass the acceptance cases in [ASSESSMENT_CHECKLIST.md](#);
11. Legal, Privacy, Security, Production and workforce representatives approve the pilot operating model, and users distinguish the five outcomes; and
12. the conformance statement identifies the level, manual controls, automated capabilities and unimplemented extensions.

9. Adoption sequence

1. Test the user task. Run the [five-case smoke pilot](#) with controlled evidence and manual assessment. Record whether users understand the answer and can resolve a hold.
2. Evaluate the operating model. Run the [comparative pilot](#), including standing approvals, evidence changes, preflight, release and closeout. Measure preparation and maintenance effort alongside decision quality.
3. Connect validated controls. Automate bounded approvals and integrations only after the required Level 3 tests pass. Independent assurance supports Level 4.

Use restricted access for a controlled-sheet pilot. Protected editing ranges allocate stewardship; confidentiality requires file or record access boundaries. Archive issued decisions immutably outside mutable intake rows or through an equivalent controlled mechanism.

Visual design, conversational interfaces, dashboards and integrations remain implementation choices. Existing 0.1 implementations follow [MIGRATION_o.2.md](#). Publication of a working draft can precede implementation assurance when its status and limitations are explicit.

CHAPTER 08 / DOCS/CONFORMANCE.MD

Conformance and validation

Version 0.2.0

All unqualified **MUST** requirements apply at every level. Manual controls can meet a requirement unless automation or machine-readable exchange is explicitly assigned to a higher level. A conformance statement identifies the organisation, profile version, jurisdictions, level, responsible authority, evidence date and limitations.

Capability	Level 1 - Recorded	Level 2 - Governed	Level 3 - Connected	Level 4 - Assured
Five dimensions, eight gates, five outcomes and gate-specific bases	Required	Required	Required	Required
Preliminary intake and decision-derived action checklist	Manual preparation allowed	Managed assignments and authoritative resolution links	Integrated assessment workflow	Independently tested
Current evidence, full authority review and prerequisite checks	Human controls allowed	Human controls allowed	Automated checks plus reserved human decisions	Independently tested
Immutable decision and reconstructable evidence versions	Controlled archive allowed	Controlled archive and standard exports	Automated snapshots and replay	Independently tested
Access control, audit history and source/derivative references	Manual controlled registers allowed	Maintained programme	Integrated registers	Independently tested
Expiry, withdrawal, incident and suspension response	Named operator and events	Managed queues and service levels	Automated impact detection and enforcement	Exercised and audited
Rule positive/boundary regression cases	Documented review	Documented review	Executable conformance suite	Independently sampled
Standing approvals, verified cohorts and exception checks	Optional; controls required if used	Maintained programme	Automated bounded matching	Effectiveness measured
Published JSON export and semantic validation	Optional	Required	Required	Required
Deterministic rule evaluation and tuple-level coverage proof	Human findings preserved	Human findings preserved	Required automation	Independently tested
Independent control assurance and comparative performance measurement	Pilot encouraged	Pilot encouraged	Pilot required before production automation	Required periodically

An implementation claiming Level 1 still needs the rights, privacy, security, validity and stop-use controls. A spreadsheet format alone establishes no conformance.

Validation layers

1. Structural: JSON Schema checks types, mandatory properties, gate cardinality and effect-specific requirements.

2. Semantic: application checks cross-record references, exact versions and record digests, time, outcome aggregation, basis kinds, rule scope, conditions and approvals.
3. Domain: qualified reviewers verify the source facts, authority completeness, interpretations, technical settings and effectiveness of duties.
4. Operational: an operator or integration confirms the approved scope and prerequisites at execution, tracks changes and evidences completion.

The first two layers cannot establish the last two. A valid export can faithfully represent an incorrect human interpretation.

Repository implementation status

`reference/semantics.js` contains executable primitives for three-valued set and retention matching, bound source rows, production-relationship derivation, gate aggregation and preflight readiness.

`reference/validation.js` validates structural schemas and selected cross-record invariants and replays frozen manual findings. The CLI validates every published JSON example.

The helpers do not implement automated legal interpretation, a registry service, full authority-set discovery, exhaustive coverage proof checking, exception authorisation, identity verification, live environment monitoring, release-duty checks or operational enforcement. The bundle demonstrates a manual/hybrid decision core. The repository itself claims no operational conformance level. A complete implementation demonstrates its selected level through the controls above.

The [assessment checklist](#) defines a derived output format, generation rules, acceptance cases and a worked example. Automated generation remains unimplemented. Existing tests cover the published schemas and reference helpers; checklist acceptance cases still require implementation validation.

CHAPTER 09 / DOCS/SCENARIOS.MD

Worked scenarios

Version 0.2.0

Fictional studio

Northstar Animation is a fictional studio with teams in Sydney, Brisbane, Los Angeles and Vancouver. It produces client-funded series, company-owned originals, co-productions and internal media. Studio, production, contributor, agreement and decision details below are fictional; references to existing products are illustrative.

These scenarios demonstrate reasoning structure rather than real legal outcomes. The same facts can produce different decisions under another contract, jurisdiction or policy. Scenario 1 has a complete machine-readable bundle; the other scenarios summarise selected facts and assume verified findings for every remaining gate. They are not complete decision records.

Scenario map

Scenario	Source	Operation	Purpose	Environment	Outcome
1. Client-show voices for an original	Completed client voice recordings	Train lip-sync model	Different, company-owned show	Private research tenant	Hold - approval required
2. Same-show meeting summary	Production-only transcript	Summarise	Same production	Approved enterprise workspace	Permitted with conditions
3. Original-show concept exploration	Cleared internal concept pack	Generate and transform	Same production	Private creative tenant	Permitted with conditions
4. Employee-message ranking	Internal email and chat	Rank people	Employment decision support	Studio-hosted model	Prohibited under current rules
5. Search across client notes	Multi-show notes and client feedback	Embed and retrieve	Reusable studio knowledge base	Private search environment	Hold - information required
6. Client-funded improvements to a studio rig	Rig code with mixed lineage	Fine-tune rigging assistant	Reusable studio capability	Studio-hosted model	Hold - information required
7. Approved source in a provider-training account	Cleared original-show art	Generate	Same production	Vendor account with training enabled	Prohibited under current rules
8. Internal-show performer experiment	Named performer recordings	Train mouth-shape predictor	Same production research	Australia-only research tenant	Permitted with conditions

1. Can we use voice recordings from a completed client show to train a lip-sync model for our own IP show?

The [assessment checklist example](#) shows how this use case becomes an initial questionnaire and an evidence-backed list of assigned questions and approval requests.

Confirmed proposal

- Source: Session HH-204 from the completed client service production Harbor Heroes. The manifest contains 12 WAV takes from Ava Lin and Mateo Ruiz.
- Legal authority and contributors: StreamWave Kids controls archived production material. Ava's rider requires a purpose-specific training agreement. Mateo's rider prohibits synthetic voice substitution and does not expressly address mouth-shape training. Different fictional collective terms apply to each performer.
- AI operation: Train and retain a model that predicts mouth shapes from dialogue audio. The technical description says the model does not generate audio, but its learned representation and deletion behaviour still require confirmation.
- Target use: Use the model on Moss & Moon, a different production owned by Northstar.
- Environment: MotionMap Lab, an Australia-hosted private tenant with provider training disabled and 30-day service-output retention. The exported model is retained in the studio research registry, where its retention duration remains unresolved.

Gate result

Gate	Result	Reason
Inventory	Satisfied	Exact takes and represented performers are known.
Source authority	Approval required	No verified client permission covers cross-production model training.
Contributor rights	Approval required	Ava requires a new agreement; Mateo's terms need an authorised interpretation and potentially new consent.
Privacy and people impact	Approval required	Purpose-specific handling and contributor expectations require review.
Operation	Information required	The retained capability, retention and deletion response need verified technical facts before interpretation.
Purpose and destination	Approval required	The proposed use crosses from client work to company IP.
Environment	Satisfied	The private tenant configuration is verified. Source permission and retained-model controls remain unresolved in other gates.
Obligations	Approval required	Consent, compensation, collective minimums, model expiry and withdrawal response are unresolved.

Decision

Hold - approval required. Do not train the model yet.

The next action is a coordinated review owned by Business and Legal Affairs. It requires:

1. written client authority covering the exact recordings, learned artefact, destination production and retention;
2. a separate conclusion for each performer under their agreement and collective terms;
3. a technical statement describing inputs, outputs, learned capability, memorisation testing, access and deletion;
4. purpose-specific agreement, compensation and withdrawal terms where required;
5. an environment approval bound to the final scope.

Client approval alone would leave the contributor gates unresolved. Performer consent alone would leave the client and cross-production gates unresolved. The collection can be narrowed and reassessed if one performer's takes are excluded.

2. Can we summarise a production meeting for the same client show?

Confirmed proposal

- Source: Transcript TT-0912 containing delivery and scheduling discussion for Tiny Titans. Personal discussion and performance commentary were removed. Participants received notice.
- Legal authority and contributors: The client agreement permits same-show workflow assistance. A standing privacy and information-governance approval covers production-only meetings.
- AI operation: Create a summary and action list. No embeddings or model training.
- Target use: Share with the authorised Tiny Titans production team.
- Environment: NotePilot Enterprise in the show's tenant; provider training disabled; source deleted within 90 days.

Decision

Permitted with conditions. The user records the meeting ID, checks the attendee and content exclusions, keeps the summary within the show team, verifies human accuracy and follows the 90-day deletion rule. A meeting containing HR, health or performance discussion requires a fresh assessment.

3. Can we use our original-show artwork for concept exploration?

Confirmed proposal

The source is a manifest of 18 internally commissioned Moss & Moon designs. Artist riders cover same-show generative exploration. Licensed references and performer likenesses are excluded. The operation generates exploratory images for the same show in a private Australian tenant with provider training disabled.

Decision

Permitted with conditions. Outputs remain within Moss & Moon, carry source-pack lineage, receive art-director review and are checked for unintended third-party or likeness content. Uploaded inputs are deleted within 14 days. Adding licensed reference images creates a new collection version and assessment.

4. Can we analyse staff messages to rank individual performance?

Confirmed proposal

The sources are employee email, chat and production metrics. The operation scores and ranks named people for management use. The environment is studio-hosted.

Decision

Prohibited under current rules. Northstar's fictional workforce AI policy expressly prohibits using communications or production metrics to score individual performance or make automated employment decisions. Hosting the model internally does not change the outcome. A separate aggregate workflow-health study could be proposed with de-identification, workforce consultation, privacy review and controls preventing individual inference.

5. Can we build a searchable AI knowledge base from all client notes?

Confirmed proposal

The initial source selection points to “all production-tracking notes” across five productions. It includes internal notes, client comments, review media links, personal names and unknown attachments. The operation creates embeddings and a persistent retrieval index for studio-wide use.

Decision

Hold - information required. The source is not a controlled collection. The data steward must inventory productions, note categories, attachments, contributors and agreements; separate client-confidential and personal content; and identify which productions permit persistent cross-show retrieval. Each approved partition receives its own manifest, access group, retention and decision. “All notes” cannot act as an assessable source.

6. Can a reusable rigging assistant learn from a rig improved on a client-funded show?

Confirmed proposal

Northstar brought a pre-existing rig framework to Forest Patrol, a client-funded service show. During production, the team added a deformation system, client character setup and show-specific scripts. The proposed operation fine-tunes a reusable rigging assistant for future productions.

Decision

Hold - information required. Ownership of the pre-existing framework does not establish rights in every improvement. Production Operations and Legal must map the repository into:

- documented Northstar background IP;
- general improvements permitted for reuse;
- client-funded project IP;
- client assets and character-specific data;
- employee and contractor contributions;
- third-party libraries and licences.

The team can create a clean collection containing only positively cleared components, then reassess training, model retention and future destinations. This scenario shows why “our rig” is an insufficient source description.

7. Can cleared artwork be used in a vendor account that improves its models from customer data?

Confirmed proposal

The source and same-show purpose are covered. The selected account allows the provider to retain inputs and outputs for service improvement and model training.

Decision

Prohibited under current rules. Northstar's active environment policy prohibits provider training on confidential production material. The requester can move the proposal to an approved private tenant with provider training disabled and run a new assessment. Source clearance and environment approval are independent gates.

8. Can we run a same-show lip-sync experiment with a performer who has agreed?

Confirmed proposal

Nia Okafor signed a fictional rider permitting named Moss & Moon takes to train a mouth-shape predictor for that production. The rider excludes voice and likeness synthesis, limits access to the Sydney and Brisbane research team, requires a research fee, requires source deletion within 30 days and sets model deletion for 31 December 2026. Applicable collective minimums have been checked. MotionMap Lab matches the stated Australian environment.

Decision

Permitted with conditions. The decision lists the exact takes, team, environment, operation and dates. It requires payment evidence, source and model deletion records, lineage and immediate stop-use review if consent is withdrawn. US access, another performer, another show or a voice-generation feature creates a new assessment.

Additional boundary cases for version 0.2

Case	Expected finding	Evidence needed
All applicable gates are satisfied, remaining gates have verified non-applicability, and no duties are outstanding	Permitted; perform a fresh preflight before processing	Complete current evidence and bounded scope
An approved internal output is proposed for public release	Reassess the distribution scope and complete release duties	Actual output, audience, channel and release-check evidence
Qualified reviewer verifies public-domain source status	Source-rights basis may be satisfied; remaining gates are assessed separately	Jurisdiction, provenance, reviewer and bounded conclusion
Licence covers the exact ordinary use	Use the licence as a rights basis; obtain extra consent only where independently required	Licence version, scope, contributor/privacy applicability findings
Qualified reviewer records an applicable legal exception	Use the scoped legal interpretation as evidence	Jurisdiction, conditions, facts and authorised reviewer; internal policy still applies
Verified material contains no identifiable people	Contributor/privacy non-applicability may be recorded where the reviewer confirms it	Verified source content, relevant rights review and named authority
Research uses another show's assets and publishes results	Preserve research, cross-production and public-release facts together	Source/destination IDs and all applicable distribution rights
Rule selects Show A and Performer B, who appears only in Show B	Rule fails to match that relationship	Bound material/person/production rows
Permit allows AU; actual processing is AU and US	Environment permission does not cover the proposal	Complete current processing-region set
Transcript retains a confidentiality restriction	Propagate the duty to the transcript	Source authority and retained applicability finding
Reviewer establishes a source-specific duty does not govern an aggregate	Record a scoped release for that output	Transformation facts and qualified interpretation; other duties remain assessed
Training-source influence cannot be removed after a withdrawal trigger	Suspend the affected model and record quarantine, retirement or retraining review	Lineage, applicability assessment, technical limits and accountable owner
Conditional approval requires a research fee before processing	Await prerequisites until payment evidence is verified	Condition event and fresh operator preflight

These cases describe evaluation structure. Positive findings require the stated evidence and all other applicable gates. They establish no general legal rule about public-domain status, exceptions, consent or derivative rights.

CHAPTER 10 / DOCS/INTEROPERABILITY.MD

Interoperability and references

Version 0.2.0

SlateCheck combines established governance and policy concepts with creative-production records. The mappings below describe conceptual relationships. They are not assertions of formal standards conformance or lossless interchange.

SlateCheck concept	Related foundation	SlateCheck-specific treatment
Permit, prohibit, conditions and constraints	ODRL permissions, prohibitions, duties and constraints	Gate-specific evidence, studio authority boundaries and contributor-level coverage
Materials, collections and governing parties	ODRL assets, asset collections and parties	Show provenance, manifests, individual riders and client-funded work
Purpose, processing, people, location and controls	Data Privacy Vocabulary (DPV)	Production relationship, retained capability, contributor and tool-environment records
Accountable ownership, review and lifecycle monitoring	NIST AI Risk Management Framework	An operational assessment, decision and obligation workflow for creative production
AI-product due diligence and ongoing review	OAIC guidance	Actual tenant, regions, provider use, retention and change-trigger evidence

ODRL can provide policy-expression mappings; DPV can provide vocabulary mappings. Implementations should publish explicit term mappings, supported subsets, extensions and information lost during conversion. For example, a SlateCheck hold for missing evidence must retain its unresolved status rather than becoming an ODRL permission. A generic purpose label must preserve SlateCheck's independent production, reuse and distribution dimensions.

A future formal ODRL profile or DPV-aligned export requires separate specification and conformance tests. SlateCheck's contribution is the integrated production-specific workflow and its evaluated usefulness. Research novelty requires a broader literature and practice comparison than these mappings alone.

Primary references

These are the editions used as design foundations. They are informative references; implementations verify the current law and guidance applicable to their own use. Links checked 16 September 2026.

- [W3C ODRL Information Model 2.2](#)
- [Data Privacy Vocabulary 2.0, W3C Community Group report](#) - a Community Group specification
- [NIST AI Risk Management Framework 1.0 \(2023\)](#)
- [OAIC guidance on commercially available AI products](#)
- [OAIC guidance on developing and training generative AI models](#)

Comparative pilot protocol

Version 0.2.0

Question and scope

Test whether SlateCheck improves decision consistency and obligation discovery while keeping preparation and maintenance effort sustainable. Begin with 15-20 controlled real cases and fictional or de-identified publication versions. Include routine approvals, mixed contributors, ordinary licensed material, missing inventory, ambiguous rights, cross-production research, public distribution, retained models and changed evidence.

Keep the pilot's jurisdiction, operations, source classes and excluded high-impact uses explicit. Existing authorised processes continue to control actual processing during evaluation. Pilot comparison decisions do not grant new permission.

Method

1. Record the existing process's outcome, reasoning, missed facts, elapsed time and specialist effort for each case. Include held, declined and abandoned cases.
2. Freeze the evidence set and record unresolved facts. Two qualified reviewers independently encode and assess each case using SlateCheck. Preserve their initial findings before discussion.
3. Compare outcomes, gate findings, obligations, next owner and explanations. A third qualified reviewer adjudicates disagreements where possible. Preserve residual uncertainty; reviewer consensus alone establishes no legal ground truth.
4. Counterbalance case order or use matched cases to limit learning effects when comparing the existing process and SlateCheck. Record reviewers' expertise and prior exposure.
5. Repeat representative uses under a bounded standing approval. Account for the initial rule-authoring and record-preparation work separately from marginal repeat-use effort.
6. Change one fact per replay: add a contributor, change a region, expire evidence, withdraw consent, change the destination, add public distribution or miss a prerequisite. Confirm the correct hold/suspension and affected derivative set.
7. Ask a reviewer outside the original pair to reconstruct a sample decision from its frozen bundle and perform a simulated preflight and closeout.
8. Follow the pilot long enough to observe at least one evidence update, repeated use and retention/obligation deadline, or explicitly simulate those events and label them as simulations.

Measures

Measure	Collection method
Outcome and gate agreement	Report raw counts and agreement rates; list unresolved cases separately
Inappropriate permission	Qualified adjudication against the frozen evidence and applicable authority
Unnecessary restriction/escalation	Review holds that could have been resolved from existing verified evidence
Missed or incorrect obligations	Compare duties, scope, owner, timing and derivative applicability
Preparation cost	Minutes spent inventorying, interpreting, authoring rules and verifying environments
Decision and maintenance cost	Specialist minutes per case, repeat use, amendment and expiry cycle
Resolution delay	Time from intake to actionable answer, separated from active staff effort
Usability and bypass	Observed comprehension, abandonment and self-reported/observed bypass attempts
Reconstruction and change response	Successful replay, affected-record discovery and suspension latency

Do not convert these measures into individual employee-performance scores. Publish denominators, case mix, incomplete cases and uncertainty. A small pilot supports feasibility conclusions; it cannot establish general legal accuracy or industry-wide effectiveness.

Go/no-go review

Before seeing results, Legal, Privacy, Security, Production and workforce representatives agree on acceptable handling of critical errors, response times and staff effort. Any unexplained permission in the presence of a verified prohibition or missing prerequisite blocks expansion until investigated. Lower error rates and operating cost must be assessed together; shorter turnaround alone is insufficient.

The review chooses a bounded expansion, another pilot iteration or suspension. Its record identifies unresolved risks, owners, evidence and the next review date. No pilot outcome may waive external rights.

CHAPTER 12 / PILOT/README.MD

Five-case usability smoke pilot

Run this smoke pilot before the larger comparative protocol in [docs/PILOT_EVALUATION.md](#). It tests whether a requester can submit a real AI-use question, understand the answer and move a held request toward resolution at reasonable cost.

SlateCheck is being evaluated. Existing authorised processes continue to control real work, and no pilot result grants permission.

Working rule: specification soft freeze

During the five cases, change the specification or prototype only to:

- correct a safety, privacy, security or factual defect;
- repair a blocking usability defect observed in a session; or
- capture information that changes an outcome, owner or next action.

Record proposed changes in [RESULTS.md](#). Defer new services, integrations, schemas, scenarios and general polish until the case-5 review.

Questions

1. Can a requester describe and confirm a proposal without specialist help?
2. Can the organisation locate and verify the evidence needed to assess it?
3. Can the requester explain the outcome, next owner and permitted next action?
4. Can at least one held request be resolved or narrowed through reassessment?
5. Is the answer worth the requester and specialist effort?

Participants

Recruit at least two requesters from production, creative or technology who did not write the framework, the evidence owners or authorities needed for the cases, and one facilitator. One person may hold several organisational roles; record each role separately. Do not use results as employee-performance data.

Select five cases

Use current real work where it can be handled safely. De-identify the pilot record and keep source media, contract text and personal information in their controlled systems.

Case	Required shape	What it tests
1	Routine same-production use with a likely standing basis	Friction on an ordinary request
2	Incomplete or poorly bounded material selection	Efficient discovery of missing inventory
3	Contributor, client or licensed-material ambiguity	Whether a hold becomes a specific request
4	Tool, account, region or retention mismatch	Whether environment controls stay independent and understandable

Case	Required shape	What it tests
5	High-impact, restricted or clearly unacceptable use	Whether the boundary and a safer alternative are understood

Include at least one case with an existing-process decision and one case the requester genuinely wants to progress. Register the five cases in `RESULTS.md` before the first session.

Run each session

Allow 30-45 minutes and use the case template in `RESULTS.md`.

1. Unassisted intake

Give the requester the prototype and this prompt:

Use SlateCheck to ask whether you can carry out this proposed AI use. Stop when you receive a decision brief.

Record time, uncertainty, guessed fields, terminology problems, sensitive-data handling, bypass attempts and facilitator interventions. Avoid coaching unless the requester becomes completely blocked.

2. Comprehension check

Hide the brief and ask:

1. What can you do now?
2. What must wait or stop?
3. Who acts next?
4. What information or approval is needed?
5. What change requires reassessment?

Rate each response against the issued brief and preserve the participant's wording when it reveals a usability problem.

3. Evidence and decision work

The relevant stewards and authorities process the case. Record active minutes for evidence location, missing-fact preparation, interpretation, data entry and decision review. Record waiting time separately and preserve uncertainty.

4. Resolution trace

For every hold, record one next action, owner, required evidence and completion test. Take at least one genuine held case from assigned work through verified response and reassessment. If it cannot progress, retain the external blocker and owner rather than simulating success.

Initial thresholds

Confirm these hypotheses with the pilot authorities before seeing results.

Expansion pauses while any of these remains unexplained:

- permission despite a verified prohibition or missing prerequisite;
- a decision that cannot be reconstructed from its stated evidence;
- an unresolved item presented as verified; or
- a hold, condition or prohibition without an accountable owner.

Usability targets are:

- median unassisted intake time of 3 minutes or less;
- in at least 4 of 5 case sessions, the requester identifies the immediate action and next owner;
- every hold identifies a specific missing fact, approval or condition;
- no requester treats saving a brief as permission; and
- the held case reaches reassessment or retains a documented external blocker.

Report actual staff minutes alongside the current-process baseline. Flag any routine case that takes materially longer without discovering a relevant fact, obligation or inconsistency.

Review after case 5

Classify proposed changes as:

- Keep - changed a decision or prevented a material error.
- Simplify - mattered but was difficult to supply or understand.
- Remove or defer - did not affect the cases and has no independent safety requirement.

Choose whether to proceed to the comparative pilot, revise and repeat the smoke pilot, or suspend implementation pending resolution of a cost, authority or evidence problem. Record the evidence, dissent, owner and next review date.

CHAPTER 13 / PILOT/RESULTS.MD

Usability smoke pilot results template

Status: Not started

Pilot owner: Unassigned

Pilot window: Unscheduled

Results review date: Unscheduled

This file is the live record for the five-case smoke pilot described in [README.md](#). Record direct observations and measured times. Label inferences, simulated events and unresolved items explicitly.

Agreed boundaries and thresholds

- Jurisdiction and locations:
- Included operations:
- Excluded high-impact uses:
- Existing process that continues to authorise real work:
- Evidence access restrictions:
- Confirmed or revised safety blockers:
- Confirmed or revised usability targets:
- Authorities agreeing thresholds and date:

Case register

ID	Case shape	De-identified question	Requester role	Evidence/decision owners	Existing-process outcome	Session date	State
USP-01	Routine same-production use						Not started
USP-02	Incomplete material selection						Not started
USP-03	Rights or authority ambiguity						Not started
USP-04	Environment mismatch						Not started
USP-05	Restricted or high-impact use						Not started

Session record template

Copy this section once for each case.

USP-XX - short case name

Prepared facts

- Question presented to requester:
- Business purpose:
- Known evidence at session start:
- Known unresolved facts at session start:
- Current-process outcome and basis, if available:
- Source references retained outside this file:

Intake observation

- Participant and role:
- Started:
- Decision brief reached:
- Unassisted intake duration:
- Facilitator interventions:
- Fields answered confidently:
- Fields guessed, misunderstood or left unknown:
- Incorrect suggestions accepted:
- Sensitive-information handling observation:
- Abandonment or bypass behaviour:
- Other direct observations:

Comprehension

Question	Participant's answer	Rating and brief basis
What can you do now?		
What must wait or stop?		
Who acts next?		
What is needed?		
What triggers reassessment?		

Decision and effort

- SlateCheck outcome:
- Current-process outcome:
- Material difference and adjudicated significance:
- Requester active minutes:
- Existing-record location/checking minutes:
- Missing evidence preparation minutes:
- Interpretation/approval minutes:
- SlateCheck data-entry minutes:

- Decision review/issue minutes:
- Elapsed waiting time:
- Relevant fact, obligation or inconsistency discovered:
- Unsafe permission, unnecessary hold or missing obligation:

Resolution trace

Date/time	Event	Owner or authority	Evidence	Result / next action
	Initial decision			

- Final or current state:
- Reassessment decision ID, if applicable:
- External blocker and owner, if still open:

Candidate changes

Observation	Direct evidence	Candidate change	Keep / simplify / remove or defer	Decision after case 5

Aggregate results

Complete after case 5. Report counts and denominators even when cases remain open.

Measure	Result	Target or comparison	Finding
Median unassisted intake time		<= 3 minutes	
Sessions with correct immediate action and next owner	/ 5	>= 4 / 5	
Holds with specific requirement and owner	/	All	
Sessions where a requester mistook a brief for permission	/ 5	0 / 5	
Unsafe permissions		0	
Unnecessary holds or escalations		Compare current process	
Missed or incorrect obligations		0 critical	
End-to-end held case reassessed		Required, or documented blocker	
Median requester active minutes		Baseline	
Median specialist active minutes		Baseline	
Median elapsed resolution time		Baseline	

After-case-5 change review

Keep

-

Simplify

-

Remove or defer

-

Go/no-go record

- Decision: Unmade
- Rationale:
- Safety blockers:
- Operating-cost finding:
- Residual uncertainty or dissent:
- Required changes:
- Owner:
- Next review date:

CHAPTER 14 / DOCS/MIGRATION_0.2.MD

Migration from 0.1 to 0.2

Version 0.2 introduces incompatible schema and semantic changes. Preserve historic 0.1 decisions and schemas with their original version identifiers. Reassessment produces new 0.2 records and successor links where the old decision was operationally issued.

Purpose

Replace the single `purpose.relationship` with:

- `productionRelationships`: derived from all source and destination production IDs;
- `businessPurposes`: production work, research/evaluation, internal operations/communications or explicit unknown;
- `reuseIntent`: task-only, production-only, reusable studio capability or unknown;
- `distribution`: internal, client delivery, external distribution, public release or unknown.

A legacy value alone cannot establish the missing dimensions. Confirm them with the requester. Research can coexist with different-production use and public release. The migrated example explicitly records research and production work on a different show with internal distribution.

Rules

Replace selector arrays with `{ "operator": "...", "values": [...] }`. Review each predicate's intended set semantics; automatic conversion to `any_of` is unsafe for permissions. Whole-proposal selectors in permit rules require `all_in` or `equals_set`. Empty selector values are invalid; omitted predicates are unrestricted within the rule's verified authority.

The old client-approval rule used alternatives in a single purpose field. Version 0.2 publishes separate atomic rules for cross-production use and reusable capability. Matching both produces one obligation per distinct required act where the verified authority and scope are identical; preserve both rule references.

Add `gates`, registry-envelope metadata, current interpretation confidence and positive/boundary test references. Requirements include timing and evidence requirements. Approval effects identify an approver and an approval requirement; information effects identify an information requirement. Numeric retention constraints belong in `constraints`, with the associated operational duty recorded separately.

Decisions and evidence

Record every core gate exactly once, plus uniquely named extension gates. Add gate-specific basis evidence, responsible roles, evaluation time and policy-profile version. Preserve every relied-on evidence and rule version in a complete snapshot. For manual/hybrid bundle replay, retain the verified gate findings as controlled evidence records and list their IDs in the evaluation method.

Collection-only intake is supported when a versioned collection and manifest are referenced. Before permission, the collection must be expanded and verified. Registry records follow the common envelope; proposed uses and decisions follow their specialised transaction schemas.

Execution and lineage

Keep execution state separate from decision outcome. Record preflight, suspension and completion events. Retain a per-obligation applicability finding for relevant derivative types. Existing derivatives with uncertain inherited duties require review of the affected use.

Verification

Run `npm run validate:spec`, `npm test` and `npm run check`. Then perform domain review and pilot validation. The repository tests check formats and selected semantics; organisational conformance requires the controls in [CONFORMANCE.md](#).

CHAPTER 15 / DOCS/PUBLICATION_REVIEW.MD

Publication review

Review date: 16 September 2026

Specification: 0.2.0 working draft

Author: Ben Gunsberger

Scope: current specification, schemas, examples, reference harness and pilot materials

Assessment

The specification is sufficiently comprehensive for external working-draft review and controlled pilot evaluation. It covers intake, individual and collective authority, technical processing, purpose and distribution, conditions, lineage, change response, appeals, privacy, conformance and evaluation. Further general chapters would add maintenance cost; expansion should follow observed pilot findings. The revised README and existing specification files total about 16,900 words, within 1% of their pre-audit length; this review is separate supporting material.

Public distribution still needs explicit reuse terms and an identified publication location and release snapshot. Operational effectiveness remains untested: the five-case pilot is not started. This review establishes editorial and technical consistency within the checks below; it supplies no legal validation or operational certification.

Findings addressed

Finding	Revision
New readers lacked a short route through the document set	Added an audience-based reading guide and clearer working-draft status; credited the author.
Normative authority and handling of conflicting prose/schema requirements were unclear	Defined document authority and the separate responsibilities of structure, semantics and conformance.
Required record fields could imply that incomplete intake must invent facts	Defined unknown, provisional and non-applicable content, including when intake stays preliminary.
Extension advice conflicted with closed exchange objects	Specified supported extension points, profile requirements and schema-version metadata.
Processing readiness did not explicitly clear distribution duties	Added a release check for actual artefacts, recipients, channels and <code>before_release</code> conditions. The harness's lack of release enforcement is explicit.
Immutable archives and deletion duties needed a shared rule	Defined retention, authorised disposal, reconstruction limits and reassessment when active evidence is lost.
Checklist example used an undefined "condition plan" kind and lacked resolution-status guidance	Made it an approval action and explained links to the authoritative resolution records.
Conformance, adoption steps and reference descriptions repeated across documents	Consolidated the conformance matrix, adoption sequence and references; shortened redundant evaluation prose.
Examples could be mistaken for complete decisions across all outcomes	Distinguished the single complete bundle from abbreviated scenarios and added unconditional-permission and release boundary cases.
Digest description omitted JavaScript serialization behaviour	Documented integer-key ordering, number serialization and Unicode preservation for independent implementations.
Pilot scoring assumed five different requesters despite requiring only two	Expressed comprehension targets and denominators per case session.

Remaining release decisions

Item	Required before public distribution
Reuse terms	Ben Gunsberger or the relevant rights holder selects terms for prose/examples and for schemas/reference code. Add the approved licence text; the current draft grants no open licence.
Canonical location and feedback	Identify the repository or publication URL and a corrections/contact route.
Release identity	Archive a dated, immutable snapshot with its tag or commit, version and change note. Keep historical releases available.
Distribution contents	Include the linked specification, schemas, fictional examples and validation harness. Publish blank pilot templates; review and de-identify any future completed results separately.

Editorial corrections can retain the working-draft version before release. After publication, issue a new identified release for changes. During o.x development, incompatible schema or decision-semantic changes require a new minor version and migration guidance; editorial corrections require an identified patch release. Record compatibility explicitly.

Validation and limits

- All five JSON examples pass structural validation; the complete bundle passes selected semantic checks and manual decision replay.
- All 43 existing tests pass, and reference JavaScript syntax checks pass.
- All 50 local Markdown links and heading anchors resolve across 17 distribution documents.
- The five primary foundation links were retrieved from their publishers. [Interoperability](#) identifies the cited editions and guidance.
- Validation environment: Node.js 25.8.1 and npm 11.11.0. Other runtime versions were not tested.

The complete published bundle ends in an approval hold. Tests exercise individual matching and execution primitives; they do not constitute a full evaluator suite for every outcome, authority combination or lifecycle event. Before production automation, implement the required cases in [Decision Specification §11](#), including release enforcement, domain validation and authority coverage. Independent ports also need digest interoperability tests.

The audit preserved existing uncommitted work and made no schema, evaluator, deployment or publication changes. Publisher decisions above remain open; no licence or external release was created.

CHAPTER 16 / DOCS/GLOSSARY.MD

Glossary

Version 0.2.0

Agreement - An authoritative contract, licence, rider, consent, collective instrument, policy, notice or approved exception.

Assessment checklist - A working view of use-case questions and actions. A preliminary checklist supports intake; an evidence-backed checklist is derived from a decision and links resolutions to authoritative records.

Basis kind - The evidence category appropriate to a gate: inventory verification, rights authority, privacy basis, technical verification, operational control or non-applicability.

Condition - A mandatory control attached to a permission, such as deletion, attribution, payment, restricted access or human review.

Contributor - A person represented in material or whose work, performance, likeness, voice, information or rights are relevant to its use.

Coverage proof - Evidence that every required material/person/operation/authority combination has an applicable basis within the assessed scope.

Decision - An immutable result for one version of a proposed use, supported by a versioned evidence snapshot.

Decision owner - The person accountable for the outcome and its review path. The owner may differ from the requester or data steward.

Derivative - Any artefact created from source material, including transcripts, summaries, labels, embeddings, indexes, outputs, datasets, adapters and model weights.

Environment - The actual tool configuration and processing boundary: provider, model, tenant, account, regions, access, retention, deletion and provider use.

Evidence - A controlled record supporting a fact or rule. Evidence has provenance, status, version, verification and access controls.

Execution state - Current operational readiness or activity, maintained separately from the historical decision outcome.

Frozen assessment bundle - Archived proposal, decision, profile, rules and controlled evidence versions sufficient to reconstruct the recorded assessment.

Lineage - The recorded relationship between source material and its derivatives.

Legal authority and contributors - The agreements, rights, people, collective terms and jurisdictions that may govern a proposed use.

Manifest - A versioned inventory of the exact items in a material collection.

Non-applicability finding - A scoped, evidenced determination by a qualified authority that a particular requirement does not apply.

Obligation applicability - A verified finding that a source duty is retained, released by authorised interpretation or uncertain for a particular derivative and transformation.

AI operation - The technical activity performed on source material and the persistent artefacts it creates. Stored as the operation domain in machine-readable records.

Policy rule - A structured, queryable interpretation with an effect, applicability scope, authority and evidence.

Positive basis - Verified evidence sufficient for the specific gate and proposed use, including rights authority, technical verification or operational controls as appropriate.

Proposed use - The versioned combination of Source, Legal authority and contributors, AI operation, Target use and Environment submitted for assessment.

Target use - The intended outcome, beneficiary, destination production, audience and business context. Stored across the purpose, production-relationship and distribution domains in machine-readable records.

Represented person - An identifiable person appearing in, speaking in, creating or otherwise linked to material in a way relevant to rights or privacy.

Source - The exact material, provenance, production context, contents and classifications involved in a proposed use.

Standing approval - A pre-authorised, bounded pattern that can be matched automatically while all facts and conditions remain valid.

Synthetic performance - Generated or transformed audio, visual or behavioural material that depicts, imitates, substitutes for or materially derives from a person's performance or identity.

Verifier - A qualified person who confirms that a structured record accurately represents its authoritative source for the stated purpose.